

"""

Crypto Service - Cryptographic operations

Provides:

- AES-256-CBC encryption via Fernet
- SHA-256 hashing for pseudonyms
- Secure key generation and management
- Evidence hash calculation
- Encryption/decryption for sensitive data

"""

```
import hashlib
import logging
import secrets
from datetime import datetime
from pathlib import Path
from typing import Optional, Tuple

from cryptography.fernet import Fernet
from cryptography.hazmat.primitives import hashes
from cryptography.hazmat.primitives.kdf.pbkdf2 import PBKDF2HMAC
```

Initialize logger

```
logger = logging.getLogger(__name__)
```

class CryptoService:

"""

Crypto Service - Cryptographic operations.

Features:

- AES-256-CBC encryption using Fernet
- SHA-256 hashing for pseudonyms and evidence
- Secure random generation
- Key derivation
- Evidence integrity verification

"""

```
def __init__(self, master_key: Optional[bytes] = None):
```

"""

Initialize crypto service.

Args:

 master_key: Master encryption key (32 bytes).

 If None, generates a new key (for testing only).

"""

```
if master_key:
    self.master_key = master_key
else:
    # Generate new key (WARNING: for testing/development only)
    self.master_key = Fernet.generate_key()
    logger.warning(
        "CryptoService initialized with auto-generated key. "
        "This should only be used for testing!"
    )

# Initialize Fernet cipher
self.cipher = Fernet(self.master_key)

logger.info("CryptoService initialized")

def generate_pseudonym(self, submission_id: str, salt: Optional[str] = None) -> str:
    """
    Generate anonymous pseudonym using SHA-256.

    Creates a collision-resistant pseudonym that preserves anonymity
    while allowing submission tracking.

    Args:
        submission_id: Unique submission identifier
        salt: Optional salt for additional randomness

    Returns:
        str: Hex-encoded pseudonym (e.g., "whistleblower-a3b5c7...")
    """
    try:
        # Create input for hash
        hash_input = submission_id.encode('utf-8')

        # Add salt if provided
        if salt:
            hash_input += salt.encode('utf-8')
        else:
            # Generate deterministic salt from submission_id
            salt_hash = hashlib.sha256(submission_id.encode()).digest()
            hash_input += salt_hash

        # Calculate SHA-256 hash
        hash_object = hashlib.sha256(hash_input)
        hash_hex = hash_object.hexdigest()

        # Create readable pseudonym (first 12 chars for readability)
        return hash_hex[:12]
    except Exception as e:
        logger.error(f"Error generating pseudonym: {e}")
        raise
```

```
pseudonym = f"whistleblower-{hash_hex[:12]}"

logger.debug(f"Generated pseudonym: {pseudonym}")

return pseudonym

except Exception as e:
    logger.error(f"Failed to generate pseudonym: {e}")
    # Fallback pseudonym
    return f"whistleblower-{secrets.token_hex(6)}"
```

def hash_data(self, data: bytes) -> str:
 """

Calculate SHA-256 hash of data.

Args:

data: Raw data bytes

Returns:

str: Hex-encoded SHA-256 hash

"""

try:

```
    hash_object = hashlib.sha256(data)
    return hash_object.hexdigest()
```

except Exception as e:

logger.error(f"Failed to hash data: {e}")

raise ValueError(f"Hash calculation failed: {str(e)}")

def hash_file(self, file_path: Path) -> str:
 """

Calculate SHA-256 hash of file.

Uses streaming to handle large files efficiently.

Args:

file_path: Path to file

Returns:

str: Hex-encoded SHA-256 hash

Raises:

FileNotFoundError: If file doesn't exist

IOError: If file read fails

"""

try:

```
if not file_path.exists():
    raise FileNotFoundError(f"File not found: {file_path}")

hash_object = hashlib.sha256()

# Stream file in chunks (64KB)
chunk_size = 65536

with open(file_path, 'rb') as f:
    while True:
        chunk = f.read(chunk_size)
        if not chunk:
            break
        hash_object.update(chunk)

file_hash = hash_object.hexdigest()

logger.debug(f"File hash: {file_hash[:16]}... (size={file_path.stat().st_size} bytes)")

return file_hash
```

```
except FileNotFoundError:
    raise
except Exception as e:
    logger.error(f"Failed to hash file {file_path}: {e}")
    raise IOError(f"File hash failed: {str(e)}")
```

def encrypt_data(self, data: bytes) -> bytes:

"""
Encrypt data using AES-256-CBC (via Fernet).

Args:

data: Raw data to encrypt

Returns:

bytes: Encrypted data (includes IV and MAC)

Raises:

ValueError: If encryption fails

"""

```
try:
    encrypted = self.cipher.encrypt(data)
    logger.debug(f"Encrypted {len(data)} bytes ! {len(encrypted)} bytes")
    return encrypted
```

```
except Exception as e:
    logger.error(f"Encryption failed: {e}")
    raise ValueError(f"Encryption failed: {str(e)}")

def decrypt_data(self, encrypted_data: bytes) -> bytes:
    """
    Decrypt data encrypted with encrypt_data().

    Args:
        encrypted_data: Encrypted data

    Returns:
        bytes: Decrypted data

    Raises:
        ValueError: If decryption fails (wrong key or corrupted data)
    """
    try:
        decrypted = self.cipher.decrypt(encrypted_data)
        logger.debug(f"Decrypted {len(encrypted_data)} bytes != {len(decrypted)} bytes")
        return decrypted
    except Exception as e:
        logger.error(f"Decryption failed: {e}")
        raise ValueError(f"Decryption failed (invalid key or corrupted data): {str(e)}")

def encrypt_file(self, input_path: Path, output_path: Path) -> str:
    """
    Encrypt file and save to output path.

    Args:
        input_path: Path to file to encrypt
        output_path: Path to save encrypted file

    Returns:
        str: SHA-256 hash of original file

    Raises:
        FileNotFoundError: If input file doesn't exist
        IOError: If encryption or write fails
    """
    try:
        if not input_path.exists():
            raise FileNotFoundError(f"Input file not found: {input_path}")
        # Read file
```

```

        with open(input_path, 'rb') as f:
            data = f.read()

        # Calculate hash before encryption
        file_hash = self.hash_data(data)

        # Encrypt
        encrypted = self.encrypt_data(data)

        # Write encrypted file
        output_path.parent.mkdir(parents=True, exist_ok=True)
        with open(output_path, 'wb') as f:
            f.write(encrypted)

        logger.info(
            f'File encrypted: {input_path.name} !' {output_path.name} "
            f"({len(data)}) !' {len(encrypted)} bytes"
        )
    )

    return file_hash

except FileNotFoundError:
    raise
except Exception as e:
    logger.error(f"File encryption failed: {e}")
    raise IOError(f"File encryption failed: {str(e)}")

```

def decrypt_file(self, input_path: Path, output_path: Path) -> None:

"""
Decrypt file and save to output path.

Args:

 input_path: Path to encrypted file
 output_path: Path to save decrypted file

Raises:

 FileNotFoundError: If input file doesn't exist
 ValueError: If decryption fails
 IOError: If write fails

"""

try:

 if not input_path.exists():
 raise FileNotFoundError(f"Input file not found: {input_path}")

 # Read encrypted file
 with open(input_path, 'rb') as f:

```
        encrypted = f.read()

    # Decrypt
    decrypted = self.decrypt_data(encrypted)

    # Write decrypted file
    output_path.parent.mkdir(parents=True, exist_ok=True)
    with open(output_path, 'wb') as f:
        f.write(decrypted)

    logger.info(
        f"File decrypted: {input_path.name} !' {output_path.name} "
        f"!'{len(encrypted)} !' {len(decrypted)} bytes"
    )

except (FileNotFoundException, ValueError):
    raise
except Exception as e:
    logger.error(f"File decryption failed: {e}")
    raise IOError(f"File decryption failed: {str(e)}")

def generate_secure_token(self, length: int = 32) -> str:
    """
    Generate cryptographically secure random token.

    Args:
        length: Token length in bytes (default: 32)

    Returns:
        str: Hex-encoded random token
    """
    return secrets.token_hex(length)

def derive_key(self, password: str, salt: bytes, iterations: int = 100000) -> bytes:
    """
    Derive encryption key from password using PBKDF2.

    Args:
        password: User password
        salt: Random salt (at least 16 bytes)
        iterations: PBKDF2 iteration count

    Returns:
        bytes: Derived key (32 bytes)
    """
    try:
```

```
        kdf = PBKDF2HMAC(  
            algorithm=hashes.SHA256(),  
            length=32,  
            salt=salt,  
            iterations=iterations  
        )  
  
        key = kdf.derive(password.encode('utf-8'))  
  
        logger.debug(f"Key derived from password (iterations={iterations})")  
  
        return key  
  
    except Exception as e:  
        logger.error(f"Key derivation failed: {e}")  
        raise ValueError(f"Key derivation failed: {str(e)}")  
  
def verify_hash(self, data: bytes, expected_hash: str) -> bool:  
    """  
    Verify data matches expected hash.  
    """
```

Args:

data: Data to verify
 expected_hash: Expected SHA-256 hash (hex)

Returns:

bool: True if hash matches, False otherwise

try:

```
    actual_hash = self.hash_data(data)  
    return secrets.compare_digest(actual_hash, expected_hash)
```

except Exception as e:

```
    logger.error(f"Hash verification failed: {e}")  
    return False
```

```
def create_submission_hash_chain(self, *data_items: bytes) -> str:  
    """
```

Create chained hash of multiple data items.

Useful for creating composite hashes of submission + evidence + metadata.

Args:

*data_items: Variable number of data items to hash

Returns:

```
    str: Hex-encoded chained hash
    """
try:
    hash_object = hashlib.sha256()

    for item in data_items:
        hash_object.update(item)

    return hash_object.hexdigest()

except Exception as e:
    logger.error(f"Chain hash failed: {e}")
    raise ValueError(f"Chain hash failed: {str(e)}")
```

```
@staticmethod
def generate_master_key() -> bytes:
    """

```

Generate new master encryption key.

WARNING: This key must be securely stored! Loss of key means loss of all encrypted data.

Returns:

```
    bytes: 32-byte master key suitable for Fernet
    """

```

```
    return Fernet.generate_key()
```

```
@staticmethod
def key_to_string(key: bytes) -> str:
    """

```

Convert key to base64 string for storage.

Args:

```
    key: Encryption key
```

Returns:

```
    str: Base64-encoded key
    """

```

```
    return key.decode('utf-8')
```

```
@staticmethod
def string_to_key(key_string: str) -> bytes:
    """

```

Convert base64 string back to key.

Args:

```
    key_string: Base64-encoded key
```

```
Returns:
```

```
    bytes: Encryption key
```

```
"""
```

```
    return key_string.encode('utf-8')
```

```
"""
```

```
Hash Chain Service - Blockchain-like immutable chain
```

```
Provides:
```

- SHA-256 hash chain for evidence custody
- Genesis block initialization
- Block addition with chaining
- Chain verification
- Tamper detection
- Proof generation

```
"""
```

```
import hashlib
import json
import logging
import threading
from datetime import datetime
from pathlib import Path
from typing import Dict, List, Optional
from filelock import FileLock
```

```
# Initialize logger
```

```
logger = logging.getLogger(__name__)
```

```
class HashChainService:
```

```
    """
```

```
    Hash Chain Service - Immutable custody chain.
```

```
    Features:
```

- SHA-256 hash chain (blockchain-like)
- Genesis block initialization
- Recursive chaining with previous hash
- Chain integrity verification
- Tamper detection
- Proof of custody generation

```
    """
```

```
def __init__(self, data_dir: Optional[Path] = None):
```

```

"""
Initialize hash chain service.

Args:
    data_dir: Base data directory (defaults to backend/data)
"""

if data_dir:
    self.data_dir = Path(data_dir)
else:
    self.data_dir = Path(__file__).parent.parent / "data"

self.chain_file = self.data_dir / "chain.json"
self.chain_lock = threading.Lock()

# Initialize chain
self._initialize_chain()

logger.info(f"HashChainService initialized (chain_file={self.chain_file})")

def _initialize_chain(self) -> None:
    """Initialize hash chain with genesis block if not exists."""
    try:
        if not self.chain_file.exists():
            logger.info("Creating genesis block...")

            genesis_block = self._create_genesis_block()
            chain_data = {
                'blocks': [genesis_block],
                'version': '1.0.0',
                'created': datetime.utcnow().isoformat()
            }

            self._save_chain(chain_data)
            logger.info(f"Genesis block created: {genesis_block['hash'][:16]}...")
    except:
        # Verify existing chain
        self.verify_chain()

    except Exception as e:
        logger.error(f"Failed to initialize chain: {e}")
        raise

def _create_genesis_block(self) -> Dict:
    """
    Create genesis block (first block in chain).

```

Returns:

 dict: Genesis block

"""

```
genesis = {
    'index': 0,
    'timestamp': datetime.utcnow().isoformat(),
    'submission_id': 'GENESIS',
    'evidence_hash': '0' * 64,
    'previous_hash': '0' * 64,
    'nonce': 0
}

# Calculate genesis hash
genesis['hash'] = self._calculate_block_hash(genesis)

return genesis
```

```
def add_block(
    self,
    submission_id: str,
    evidence_hash: str,
    metadata: Optional[Dict] = None
) -> Dict:
    """
```

Add new block to chain.

Args:

 submission_id: Unique submission identifier
 evidence_hash: SHA-256 hash of evidence
 metadata: Optional metadata to include

Returns:

 dict: New block

Raises:

 ValueError: If block addition fails

try:

```
    with self.chain_lock:
        chain_data = self.load_chain()
        blocks = chain_data.get('blocks', [])
```

 # Get previous block

 if not blocks:

```
            raise ValueError("Chain has no blocks (genesis missing)")
```

```

previous_block = blocks[-1]

# Create new block
new_block = {
    'index': len(blocks),
    'timestamp': datetime.utcnow().isoformat(),
    'submission_id': submission_id,
    'evidence_hash': evidence_hash,
    'previous_hash': previous_block['hash'],
    'metadata': metadata or {}
}

# Calculate hash
new_block['hash'] = self._calculate_block_hash(new_block)

# Add to chain
blocks.append(new_block)
chain_data['blocks'] = blocks

# Save chain
self._save_chain(chain_data)

logger.info(
    f"Added block {new_block['index']} for {submission_id} "
    f"(hash={new_block['hash'][:16]}...)"
)

return new_block

except Exception as e:
    logger.error(f"Failed to add block: {e}")
    raise ValueError(f"Failed to add block: {str(e)}")

def verify_chain(self) -> bool:
    """
    Verify integrity of entire chain.

    Returns:
        bool: True if chain is valid, False otherwise
    Raises:
        ValueError: If chain is corrupted
    """
    try:
        chain_data = self.load_chain()
        blocks = chain_data.get('blocks', [])

```

```

if not blocks:
    raise ValueError("Chain is empty")

# Verify genesis block
genesis = blocks[0]
if genesis['index'] != 0:
    raise ValueError("Invalid genesis block index")

# Verify each block
for i in range(1, len(blocks)):
    current_block = blocks[i]
    previous_block = blocks[i - 1]

    # Check index continuity
    if current_block['index'] != i:
        raise ValueError(f"Block {i} has invalid index: {current_block['index']}")

    # Check previous hash linkage
    if current_block['previous_hash'] != previous_block['hash']:
        raise ValueError(
            f"Block {i} has invalid previous_hash: "
            f"expected {previous_block['hash']}, "
            f"got {current_block['previous_hash']}"
        )

    # Verify block hash
    calculated_hash = self._calculate_block_hash(current_block)
    if current_block['hash'] != calculated_hash:
        raise ValueError(
            f"Block {i} has invalid hash: "
            f"expected {calculated_hash}, "
            f"got {current_block['hash']}"
        )

logger.debug(f"Chain verified successfully ({len(blocks)} blocks)")
return True

except ValueError as e:
    logger.error(f"Chain verification failed: {e}")
    raise
except Exception as e:
    logger.error(f"Chain verification error: {e}")
    return False

def get_proof(self, submission_id: str) -> Dict:

```

```
"""
Get chain proof for submission.

Args:
    submission_id: Submission identifier

Returns:
    dict: Chain proof with block details
"""

try:
    chain_data = self.load_chain()
    blocks = chain_data.get('blocks', [])

    # Find block for submission
    for block in blocks:
        if block.get('submission_id') == submission_id:
            return {
                'block_index': block['index'],
                'block_hash': block['hash'],
                'previous_hash': block['previous_hash'],
                'timestamp': block['timestamp'],
                'evidence_hash': block['evidence_hash'],
                'chain_length': len(blocks),
                'verified': True
            }

    logger.debug(f"No proof found for submission {submission_id}")
    return {}

except Exception as e:
    logger.error(f"Failed to get proof: {e}")
    return {}

def get_chain_statistics(self) -> Dict:
"""
Get chain statistics.

Returns:
    dict: Chain statistics
"""

try:
    chain_data = self.load_chain()
    blocks = chain_data.get('blocks', [])

    # Calculate statistics
    if blocks:
```

```

        first_block = blocks[0]
        last_block = blocks[-1]

        first_time = datetime.fromisoformat(first_block['timestamp'])
        last_time = datetime.fromisoformat(last_block['timestamp'])
        chain_age_seconds = (last_time - first_time).total_seconds()
    else:
        chain_age_seconds = 0

    stats = {
        'total_blocks': len(blocks),
        'chain_version': chain_data.get('version', '1.0.0'),
        'genesis_timestamp': blocks[0]['timestamp'] if blocks else None,
        'latest_timestamp': blocks[-1]['timestamp'] if blocks else None,
        'chain_age_seconds': chain_age_seconds,
        'chain_age_hours': round(chain_age_seconds / 3600, 2),
        'average_block_time_seconds': (
            chain_age_seconds / (len(blocks) - 1)
            if len(blocks) > 1 else 0
        ),
        'chain_verified': self.verify_chain(),
        'timestamp': datetime.utcnow().isoformat()
    }

    return stats

except Exception as e:
    logger.error(f"Failed to get chain statistics: {e}")
    return {
        'error': str(e),
        'timestamp': datetime.utcnow().isoformat()
    }

def load_chain(self) -> Dict:
    """
    Load chain data from file.

    Returns:
        dict: Chain data with blocks
    """
    try:
        if not self.chain_file.exists():
            logger.warning("Chain file not found")
            return {'blocks': []}

        lock_path = self.chain_file.with_suffix('.lock')

```

```

lock = FileLock(str(lock_path), timeout=10)

with lock:
    with open(self.chain_file, 'r', encoding='utf-8') as f:
        chain_data = json.load(f)

    return chain_data

except Exception as e:
    logger.error(f"Failed to load chain: {e}")
    return {'blocks': []}

def _save_chain(self, chain_data: Dict) -> None:
    """
    Save chain data to file with atomic write.

    Args:
        chain_data: Chain data to save
    """
    try:
        lock_path = self.chain_file.with_suffix('.lock')
        lock = FileLock(str(lock_path), timeout=10)

        with lock:
            # Write to temp file first
            temp_path = self.chain_file.with_suffix('.tmp')

            with open(temp_path, 'w', encoding='utf-8') as f:
                json.dump(chain_data, f, indent=2, ensure_ascii=False)

            # Atomic rename
            temp_path.replace(self.chain_file)

    except Exception as e:
        logger.error(f"Failed to save chain: {e}")
        raise

def _calculate_block_hash(self, block: Dict) -> str:
    """
    Calculate SHA-256 hash of block.

    Args:
        block: Block dictionary

    Returns:
        str: Hex-encoded SHA-256 hash
    """

```

```
"""
# Create deterministic string representation
# Exclude 'hash' field itself
block_copy = {k: v for k, v in block.items() if k != 'hash'}

# Sort keys for consistency
block_string = json.dumps(block_copy, sort_keys=True, ensure_ascii=False)

# Calculate SHA-256
hash_object = hashlib.sha256(block_string.encode('utf-8'))
return hash_object.hexdigest()

def export_chain(self, output_path: Path) -> bool:
"""
    Export chain to file.

    Args:
        output_path: Path to export file

    Returns:
        bool: True if successful
    """
    try:
        chain_data = self.load_chain()

        with open(output_path, 'w', encoding='utf-8') as f:
            json.dump(chain_data, f, indent=2, ensure_ascii=False)

        logger.info(f"Chain exported to {output_path}")
        return True

    except Exception as e:
        logger.error(f"Failed to export chain: {e}")
        return False

def rebuild_chain(self, backup_path: Path) -> bool:
"""
    Rebuild chain from backup.

    Args:
        backup_path: Path to backup file

    Returns:
        bool: True if successful
    """
    try:
```

```

        with open(backup_path, 'r', encoding='utf-8') as f:
            chain_data = json.load(f)

        # Verify chain before rebuilding
        blocks = chain_data.get('blocks', [])

        if not blocks:
            raise ValueError("Backup chain is empty")

        # Save backup chain
        self._save_chain(chain_data)

        # Verify rebuilt chain
        self.verify_chain()

        logger.info(f"Chain rebuilt from {backup_path}")
        return True

    except Exception as e:
        logger.error(f"Failed to rebuild chain: {e}")
        return False

```

"""

Metadata Service - File metadata removal

Provides:

- EXIF data stripping from images
- GPS location removal
- Camera/device information removal
- Timestamp sanitization
- Metadata anonymization

"""

```

import logging
from datetime import datetime
from pathlib import Path
from typing import Dict, Optional

from PIL import Image
from PIL.ExifTags import TAGS

# Initialize logger
logger = logging.getLogger(__name__)

```

class MetadataService:

"""

Metadata Service - Remove sensitive metadata from files.

Features:

- EXIF data removal from images (JPEG, PNG, TIFF, etc.)
- GPS location stripping
- Camera model/serial number removal
- Timestamp anonymization
- File metadata sanitization

Protects whistleblower anonymity by removing identifying information.

"""

```
def __init__(self):
    """Initialize metadata service."""
    logger.info("MetadataService initialized")
```

```
def strip_image_metadata(
    self,
    input_path: Path,
    output_path: Path,
    preserve_orientation: bool = True
) -> Dict:
    """
```

Strip all EXIF/metadata from image file.

Args:

```
    input_path: Path to input image
    output_path: Path to save cleaned image
    preserve_orientation: Keep image orientation (recommended)
```

Returns:

```
    dict: Report of stripped metadata
```

Raises:

```
    FileNotFoundError: If input file doesn't exist
    ValueError: If file is not a valid image
    IOError: If processing fails
    """
```

try:

```
    if not input_path.exists():
        raise FileNotFoundError(f"Input file not found: {input_path}")
```

```
    # Open image
```

```
    with Image.open(input_path) as img:
```

```
        # Extract metadata before stripping (for logging)
```

```

original_metadata = self._extract_metadata(img)

# Get image format
img_format = img.format or 'JPEG'

# Handle orientation if needed
if preserve_orientation and hasattr(img, '_getexif'):
    img = self._correct_orientation(img)

# Create new image without metadata
# Convert to RGB if necessary (removes alpha channel issues)
if img.mode not in ('RGB', 'L'):
    if img.mode == 'RGBA':
        # Create white background
        background = Image.new('RGB', img.size, (255, 255, 255))
        background.paste(img, mask=img.split()[3]) # Use alpha as mask
        img = background
    else:
        img = img.convert('RGB')

# Save without metadata
output_path.parent.mkdir(parents=True, exist_ok=True)

# Save parameters (no EXIF)
save_kwargs = {
    'format': img_format,
    'quality': 95, # High quality
    'optimize': True
}

# For JPEG, explicitly disable EXIF
if img_format.upper() in ('JPEG', 'JPG'):
    save_kwargs['exif'] = b""

img.save(output_path, **save_kwargs)

# Verify metadata was removed
cleaned_metadata = self._get_metadata_summary(output_path)

report = {
    'input_file': str(input_path),
    'output_file': str(output_path),
    'original_metadata_count': len(original_metadata),
    'cleaned_metadata_count': len(cleaned_metadata),
    'metadata_removed': original_metadata,
    'sensitive_data_stripped': self._identify_sensitive_fields(original_metadata),
}

```

```
        'timestamp': datetime.utcnow().isoformat()
    }

    logger.info(
        f"Metadata stripped: {input_path.name} !' {output_path.name} "
        f"({len(original_metadata)}) fields removed"
    )

    return report

except FileNotFoundError:
    raise
except Exception as e:
    logger.error(f"Metadata stripping failed for {input_path}: {e}")
    raise IOError(f"Metadata stripping failed: {str(e)}")

def _extract_metadata(self, img: Image.Image) -> Dict:
    """
    Extract all EXIF/metadata from image.

    Args:
        img: PIL Image object

    Returns:
        dict: Metadata key-value pairs
    """
    metadata = {}

    try:
        # Try to get EXIF data
        exif_data = img._getexif()

        if exif_data:
            for tag_id, value in exif_data.items():
                # Get human-readable tag name
                tag = TAGS.get(tag_id, tag_id)

                # Convert bytes to string for logging
                if isinstance(value, bytes):
                    try:
                        value = value.decode('utf-8', errors='ignore')
                    except:
                        value = f"<binary data: {len(value)} bytes>"

                metadata[tag] = str(value)

    except:
        logger.error(f"Metadata stripping failed for {input_path}: {e}")
        raise IOError(f"Metadata stripping failed: {str(e)}")
```

```
except AttributeError:
    # No EXIF data
    pass
except Exception as e:
    logger.warning(f"Failed to extract metadata: {e}")

return metadata

def _correct_orientation(self, img: Image.Image) -> Image.Image:
    """
    Correct image orientation based on EXIF data.

    Args:
        img: PIL Image object

    Returns:
        Image: Oriented image
    """
    try:
        exif = img._getexif()

        if exif:
            # Orientation tag
            orientation_tag = 0x0112

            if orientation_tag in exif:
                orientation = exif[orientation_tag]

                # Rotation mappings
                rotations = {
                    3: 180,
                    6: 270,
                    8: 90
                }

                if orientation in rotations:
                    img = img.rotate(rotations[orientation], expand=True)
                    logger.debug(f"Corrected orientation: {orientation}")

    except Exception as e:
        logger.warning(f"Orientation correction failed: {e}")

    return img

def _get_metadata_summary(self, image_path: Path) -> Dict:
    """
```

Get summary of remaining metadata in image.

Args:

 image_path: Path to image

Returns:

 dict: Metadata summary

"""

try:

 with Image.open(image_path) as img:

 return self._extract_metadata(img)

except Exception as e:

 logger.warning(f"Failed to get metadata summary: {e}")

 return {}

def _identify_sensitive_fields(self, metadata: Dict) -> list:

"""

Identify sensitive metadata fields that were removed.

Args:

 metadata: Metadata dictionary

Returns:

 list: List of sensitive field names

"""

sensitive_fields = []

Common sensitive EXIF tags

sensitive_tags = [
 'GPSInfo', # GPS location
 'GPSLatitude',
 'GPSLongitude',
 'GPSAltitude',
 'GPSTimeStamp',
 'Make', # Camera manufacturer
 'Model', # Camera model
 'Software', # Editing software
 'DateTime', # Capture time
 'DateTimeOriginal',
 'DateTimeDigitized',
 'HostComputer', # Computer name
 'Copyright', # Copyright info
 'Artist', # Artist/creator
 'CameraSerialNumber',
 'LensSerialNumber',
 'OwnerName',


```
        'format': img.format,
        'size': img.size,
        'mode': img.mode,
        'metadata_count': len(metadata),
        'has_gps': any(k.startswith('GPS') for k in metadata.keys()),
        'has_camera_info': 'Make' in metadata or 'Model' in metadata,
        'has_timestamp': any(
            k in metadata for k in ['DateTime', 'DateTimeOriginal']
        ),
        'sensitive_fields': self._identify_sensitive_fields(metadata),
        'all_metadata': metadata
    }

    return report
```

except Exception as e:

```
    logger.error(f"Failed to get metadata report: {e}")
    return {
        'error': str(e),
        'file_path': str(image_path)
    }
```

def sanitize_filename(self, filename: str) -> str:

"""

Sanitize filename to remove identifying information.

Args:

filename: Original filename

Returns:

str: Sanitized filename

"""

Get extension

path = Path(filename)

extension = path.suffix

Generate anonymous name

timestamp = datetime.utcnow().strftime('%Y%m%d_%H%M%S')

sanitized = f"evidence_{timestamp}{extension}"

```
logger.debug(f"Filename sanitized: {filename} ! {sanitized}")
```

return sanitized

def batch_strip_metadata(

self,

```
    input_dir: Path,
    output_dir: Path,
    extensions: Optional[list] = None
) -> Dict:
"""
    Strip metadata from all images in directory.

Args:
    input_dir: Input directory
    output_dir: Output directory
    extensions: Image extensions to process (default: jpg, png, jpeg)

Returns:
    dict: Batch processing report
"""
if extensions is None:
    extensions = ['.jpg', '.jpeg', '.png', '.tiff', '.bmp']

processed = []
failed = []

try:
    output_dir.mkdir(parents=True, exist_ok=True)

    # Find all images
    image_files = []
    for ext in extensions:
        image_files.extend(input_dir.glob(f"*{ext}"))
        image_files.extend(input_dir.glob(f"*{ext.upper()}"))

    logger.info(f"Found {len(image_files)} images to process")

    # Process each image
    for input_path in image_files:
        try:
            output_path = output_dir / input_path.name

            report = self.strip_image_metadata(input_path, output_path)
            processed.append(report)

        except Exception as e:
            logger.error(f"Failed to process {input_path}: {e}")
            failed.append({
                'file': str(input_path),
                'error': str(e)
            })


```

```

batch_report = {
    'input_directory': str(input_dir),
    'output_directory': str(output_dir),
    'total_files': len(image_files),
    'processed': len(processed),
    'failed': len(failed),
    'processed_files': processed,
    'failed_files': failed,
    'timestamp': datetime.utcnow().isoformat()
}

logger.info(
    f"Batch processing complete: {len(processed)} processed, "
    f"{len(failed)} failed"
)

return batch_report

except Exception as e:
    logger.error(f"Batch processing failed: {e}")
    return {
        'error': str(e),
        'processed': len(processed),
        'failed': len(failed)
    }

```

"""

Metrics Service - Performance metrics tracking with Prometheus integration

Provides:

- Request latency tracking
- Throughput monitoring
- Memory usage tracking
- Model inference time tracking
- System resource monitoring
- Metrics persistence
- Prometheus metrics export

"""

```

import logging
import os
import psutil
import time
from collections import defaultdict, deque
from datetime import datetime, timedelta

```

```

from pathlib import Path
from typing import Dict, List, Optional, Any

# Prometheus imports (optional - graceful degradation if not available)
try:
    from prometheus_client import Counter, Gauge, Histogram
    PROMETHEUS_AVAILABLE = True
except ImportError:
    PROMETHEUS_AVAILABLE = False
    Counter = Gauge = Histogram = None

# Initialize logger
logger = logging.getLogger(__name__)

```

class MetricsService:

"""

Metrics Service - Performance and system metrics tracking with Prometheus integration.

Features:

- Request latency tracking (percentiles)
 - Throughput monitoring (requests per time window)
 - Memory usage tracking
 - Model inference time tracking
 - Error rate monitoring
 - Metrics aggregation and persistence
 - Prometheus metrics export (if available)
- """

```

def __init__(
    self,
    metrics_file: Optional[Path] = None,
    window_size: int = 1000,
    persist_interval: int = 300,
    prometheus_metrics: Optional[Dict[str, Any]] = None
):
    """

```

Initialize metrics service.

Args:

```

    metrics_file: Path to metrics persistence file
    window_size: Number of recent metrics to keep in memory
    persist_interval: Persist metrics every N seconds
    prometheus_metrics: Dict of Prometheus metric objects from main.py
"""

```

```

self.metrics_file = metrics_file
self.window_size = window_size
self.persist_interval = persist_interval

# Prometheus metrics (passed from main.py)
self.prometheus_metrics = prometheus_metrics or {}
self.prometheus_enabled = PROMETHEUS_AVAILABLE and
bool(prometheus_metrics)

if self.prometheus_enabled:
    logger.info("Prometheus metrics integration enabled")
else:
    logger.info("Prometheus metrics not available - using local metrics only")

# Request metrics (sliding window)
self.request_latencies: deque = deque(maxlen=window_size)
self.request_timestamps: deque = deque(maxlen=window_size)

# Model metrics
self.model_inference_times: Dict[str, deque] = defaultdict(
    lambda: deque(maxlen=window_size)
)

# Endpoint metrics
self.endpoint_counts: Dict[str, int] = defaultdict(int)
self.endpoint_errors: Dict[str, int] = defaultdict(int)

# Submission metrics
self.submission_stats = {
    'total': 0,
    'pending': 0,
    'completed': 0,
    'failed': 0
}

# Layer-specific metrics
self.layer_metrics = {
    'layer1_anonymity_violations': 0,
    'layer2_credibility_scores': deque(maxlen=100),
    'layer2_deepfakes_detected': 0,
    'layer3_coordination_detected': 0,
    'layer4_consensus_iterations': deque(maxlen=100),
    'layer4_convergence_times': deque(maxlen=100),
    'layer5_counter_evidence_count': 0,
    'layer5_impacts': deque(maxlen=100),
    'layer6_reports_generated': 0
}

```

```

}

# Security metrics
self.security_metrics = {
    'hash_chain_failures': 0,
    'validation_failures': defaultdict(int),
    'crypto_failures': defaultdict(int),
    'rate_limit_blocks': defaultdict(int)
}

# Storage metrics
self.storage_metrics = {
    'corruption_detected': 0,
    'index_inconsistencies': 0,
    'backup_failures': 0
}

# Queue metrics
self.queue_metrics = {
    'pending_jobs': 0,
    'processing_jobs': 0,
    'job_failures': defaultdict(int)
}

# Research/evaluation metrics
self.evaluation_metrics = {
    'auroc_score': 0.0,
    'precision_score': 0.0,
    'recall_score': 0.0
}

# System metrics
self.process = psutil.Process(os.getpid())
self.generic_gauges: Dict[str, float] = {}
self.start_time = time.time()
self.last_persist_time = time.time()

# Counters
self.total_requests = 0
self.total_errors = 0

logger.info("MetricsService initialized")

# ===== REQUEST METRICS =====

def record_request(

```

```
    self,
    endpoint: str,
    latency: float,
    status_code: int
) -> None:
"""
    Record HTTP request metrics.

Args:
    endpoint: API endpoint path
    latency: Request latency in seconds
    status_code: HTTP status code
"""
try:
    # Record latency
    self.request_latencies.append(latency)
    self.request_timestamps.append(time.time())

    # Record endpoint
    self.endpoint_counts[endpoint] += 1

    # Record errors (4xx, 5xx)
    if status_code >= 400:
        self.endpoint_errors[endpoint] += 1
        self.total_errors += 1

    self.total_requests += 1

    # Check if should persist
    if time.time() - self.last_persist_time > self.persist_interval:
        self._persist_metrics()

except Exception as e:
    logger.error(f"Failed to record request metric: {e}")
```

```
def record_model_inference(
    self,
    model_name: str,
    inference_time: float
) -> None:
"""
    Record model inference time.

Args:
```

```
    model_name: Name of ML model
    inference_time: Inference time in seconds
```

```

"""
try:
    self.model_inference_times[model_name].append(inference_time)

    logger.debug(
        f"Model inference: {model_name} "
        f"({inference_time:.3f}s)"
    )
except Exception as e:
    logger.error(f"Failed to record model metric: {e}")

def record_submission(self, status: str) -> None:
    """
    Record submission status change.

    Args:
        status: Submission status (pending/completed/failed)
    """
    try:
        self.submission_stats['total'] += 1

        if status in self.submission_stats:
            self.submission_stats[status] += 1

        # Update Prometheus metric
        if self.prometheus_enabled and 'submission_total' in self.prometheus_metrics:
            self.prometheus_metrics['submission_total'].labels(status=status).inc()

    except Exception as e:
        logger.error(f"Failed to record submission metric: {e}")

# ===== LAYER-SPECIFIC METRICS =====

def record_anonymityViolation(self) -> None:
    """
    Record Layer 1 anonymity violation.
    """
    try:
        self.layer_metrics['layer1_anonymity_violations'] += 1
        logger.warning("Anonymity violation detected")
    except Exception as e:
        logger.error(f"Failed to record anonymity violation: {e}")

def update_credibility_score(self, score: float) -> None:
    """
    Update Layer 2 credibility score.

```

```

Args:
    score: Credibility score (0.0 to 1.0)
    """
try:
    self.layer_metrics['layer2_credibility_scores'].append(score)

    # Update Prometheus gauge with average
    if self.prometheus_enabled and 'layer2_score' in self.prometheus_metrics:
        scores = self.layer_metrics['layer2_credibility_scores']
        avg_score = sum(scores) / len(scores) if scores else 0.0
        self.prometheus_metrics['layer2_score'].set(avg_score)

    logger.debug(f"Credibility score updated: {score:.3f}")

except Exception as e:
    logger.error(f"Failed to update credibility score: {e}")

def record_deepfake_detection(self) -> None:
    """Record Layer 2 deepfake detection."""
    try:
        self.layer_metrics['layer2_deepfakes_detected'] += 1
        logger.info("Deepfake detected")
    except Exception as e:
        logger.error(f"Failed to record deepfake detection: {e}")

def record_coordination_detection(self) -> None:
    """Record Layer 3 coordination detection."""
    try:
        self.layer_metrics['layer3_coordination_detected'] += 1

        # Update Prometheus counter
        if self.prometheus_enabled and 'layer3_coordination' in
self.prometheus_metrics:
            self.prometheus_metrics['layer3_coordination'].inc()

        logger.warning("Coordinated attack detected")

    except Exception as e:
        logger.error(f"Failed to record coordination detection: {e}")

def update_consensus_metrics(
    self,
    iterations: int,
    convergence_time: float
) -> None:
    """

```

Update Layer 4 consensus metrics.

Args:

```
    iterations: Number of consensus iterations
    convergence_time: Time to reach consensus (seconds)
```

"""

try:

```
    self.layer_metrics['layer4_consensus_iterations'].append(iterations)
    self.layer_metrics['layer4_convergence_times'].append(convergence_time)
```

Update Prometheus gauges

if self.prometheus_enabled:

```
    if 'layer4_iterations' in self.prometheus_metrics:
        iter_list = self.layer_metrics['layer4_consensus_iterations']
        avg_iterations = sum(iter_list) / len(iter_list) if iter_list else 0.0
        self.prometheus_metrics['layer4_iterations'].set(avg_iterations)
```

if 'layer4_convergence' in self.prometheus_metrics:

```
    self.prometheus_metrics['layer4_convergence'].set(convergence_time)
```

logger.debug(

```
    f"Consensus metrics: {iterations} iterations, "
    f"{convergence_time:.3f}s convergence time"
)
```

except Exception as e:

```
    logger.error(f"Failed to update consensus metrics: {e}")
```

def record_counter_evidence(self, impact_percent: float) -> None:

"""

Record Layer 5 counter-evidence submission.

Args:

```
    impact_percent: Impact percentage on credibility score
```

"""

try:

```
    self.layer_metrics['layer5_counter_evidence_count'] += 1
    self.layer_metrics['layer5_impacts'].append(impact_percent)
```

Update Prometheus metrics

if self.prometheus_enabled:

```
    if 'layer5_counter' in self.prometheus_metrics:
        self.prometheus_metrics['layer5_counter'].inc()
```

if 'layer5_impact' in self.prometheus_metrics:

```
    impacts = self.layer_metrics['layer5_impacts']
```

```

        avg_impact = sum(impacts) / len(impacts) if impacts else 0.0
        self.prometheus_metrics['layer5_impact'].set(avg_impact)

    logger.info(f"Counter-evidence recorded: {impact_percent:.1f}% impact")

except Exception as e:
    logger.error(f"Failed to record counter-evidence: {e}")

def record_report_generated(self) -> None:
    """Record Layer 6 report generation."""
    try:
        self.layer_metrics['layer6_reports_generated'] += 1
        logger.info("Forensic report generated")
    except Exception as e:
        logger.error(f"Failed to record report generation: {e}")

# ===== SECURITY METRICS =====

def record_hash_chain_failure(self) -> None:
    """Record hash chain validation failure."""
    try:
        self.security_metrics['hash_chain_failures'] += 1
        logger.error("Hash chain validation failure")
    except Exception as e:
        logger.error(f"Failed to record hash chain failure: {e}")

def record_validation_failure(self, validation_type: str) -> None:
    """
    Record input validation failure.

    Args:
        validation_type: Type of validation that failed
    """
    try:
        self.security_metrics['validation_failures'][validation_type] += 1
        logger.warning(f"Validation failure: {validation_type}")
    except Exception as e:
        logger.error(f"Failed to record validation failure: {e}")

def record_crypto_failure(self, operation: str) -> None:
    """
    Record cryptographic operation failure.

    Args:
        operation: Crypto operation that failed
    """

```

```
try:
    self.security_metrics['crypto_failures'][operation] += 1
    logger.error(f"Cryptographic failure: {operation}")
except Exception as e:
    logger.error(f"Failed to record crypto failure: {e}")
```

```
def record_rate_limit_block(self, reason: str) -> None:
```

```
"""
```

```
    Record rate limiter block.
```

```
Args:
```

```
    reason: Reason for block
```

```
"""
```

```
try:
```

```
    self.security_metrics['rate_limit_blocks'][reason] += 1
    logger.warning(f"Rate limit block: {reason}")
```

```
except Exception as e:
```

```
    logger.error(f"Failed to record rate limit block: {e}")
```

```
# ====== STORAGE METRICS ======
```

```
def record_storage_corruption(self) -> None:
```

```
    """Record storage corruption detection."""
```

```
try:
```

```
    self.storage_metrics['corruption_detected'] += 1
    logger.error("Storage corruption detected")
```

```
except Exception as e:
```

```
    logger.error(f"Failed to record storage corruption: {e}")
```

```
def record_index_inconsistency(self) -> None:
```

```
    """Record storage index inconsistency."""
```

```
try:
```

```
    self.storage_metrics['index_inconsistencies'] += 1
    logger.warning("Storage index inconsistency detected")
```

```
except Exception as e:
```

```
    logger.error(f"Failed to record index inconsistency: {e}")
```

```
def record_backup_failure(self) -> None:
```

```
    """Record backup failure."""
```

```
try:
```

```
    self.storage_metrics['backup_failures'] += 1
    logger.error("Backup failure")
```

```
except Exception as e:
```

```
    logger.error(f"Failed to record backup failure: {e}")
```

```
# ====== QUEUE METRICS ======
```

```

def update_queue_metrics(
    self,
    pending: int,
    processing: int
) -> None:
    """
    Update queue metrics.

    Args:
        pending: Number of pending jobs
        processing: Number of processing jobs
    """
    try:
        self.queue_metrics['pending_jobs'] = pending
        self.queue_metrics['processing_jobs'] = processing

        # Update Prometheus gauges
        if self.prometheus_enabled:
            if 'queue_pending' in self.prometheus_metrics:
                self.prometheus_metrics['queue_pending'].set(pending)
            if 'queue_processing' in self.prometheus_metrics:
                self.prometheus_metrics['queue_processing'].set(processing)

    except Exception as e:
        logger.error(f"Failed to update queue metrics: {e}")

def record_queue_failure(self, job_type: str) -> None:
    """
    Record queue job failure.

    Args:
        job_type: Type of job that failed
    """
    try:
        self.queue_metrics['job_failures'][job_type] += 1
        logger.error(f"Queue job failure: {job_type}")
    except Exception as e:
        logger.error(f"Failed to record queue failure: {e}")

# ===== EVALUATION METRICS =====

def update_evaluation_scores(
    self,
    auroc: Optional[float] = None,
    precision: Optional[float] = None,

```

```

    recall: Optional[float] = None
) -> None:
"""
    Update research evaluation scores.

Args:
    auroc: AUROC score
    precision: Precision score
    recall: Recall score
"""
try:
    if auroc is not None:
        self.evaluation_metrics['auroc_score'] = auroc
    if precision is not None:
        self.evaluation_metrics['precision_score'] = precision
    if recall is not None:
        self.evaluation_metrics['recall_score'] = recall

    logger.info(
        f"Evaluation scores updated: "
        f"AUROC={auroc}, Precision={precision}, Recall={recall}"
    )
except Exception as e:
    logger.error(f"Failed to update evaluation scores: {e}")

# ===== METRICS RETRIEVAL =====

def get_request_stats(self) -> Dict:
"""
    Get request statistics.

Returns:
    dict: Request metrics
"""
if not self.request_latencies:
    return {
        'total_requests': 0,
        'avg_latency': 0.0,
        'p50_latency': 0.0,
        'p95_latency': 0.0,
        'p99_latency': 0.0
    }

# Calculate percentiles
sorted_latencies = sorted(self.request_latencies)

```

```

n = len(sorted_latencies)

p50_idx = int(n * 0.50)
p95_idx = int(n * 0.95)
p99_idx = int(n * 0.99)

stats = {
    'total_requests': self.total_requests,
    'total_errors': self.total_errors,
    'error_rate': round(self.total_errors / self.total_requests * 100, 2) if
self.total_requests > 0 else 0.0,
    'recent_requests': len(self.request_latencies),
    'avg_latency': round(sum(self.request_latencies) / n, 3),
    'min_latency': round(min(self.request_latencies), 3),
    'max_latency': round(max(self.request_latencies), 3),
    'p50_latency': round(sorted_latencies[p50_idx], 3),
    'p95_latency': round(sorted_latencies[p95_idx], 3),
    'p99_latency': round(sorted_latencies[p99_idx], 3)
}

return stats

```

```
def get_throughput(self, window_seconds: int = 60) -> float:
```

```
"""
```

```
Calculate request throughput (requests per second).
```

Args:

 window_seconds: Time window in seconds

Returns:

 float: Requests per second

```
"""
```

if not self.request_timestamps:

 return 0.0

```
current_time = time.time()
```

```
cutoff_time = current_time - window_seconds
```

Count requests in window

```
recent_requests = sum(
```

 1 for ts in self.request_timestamps

 if ts > cutoff_time

```
)
```

Calculate throughput

```
throughput = recent_requests / window_seconds
```

```
    return round(throughput, 2)

def get_model_stats(self) -> Dict:
    """
    Get model inference statistics.

    Returns:
        dict: Model metrics by model name
    """
    model_stats = {}

    for model_name, times in self.model_inference_times.items():
        if times:
            model_stats[model_name] = {
                'total_inferences': len(times),
                'avg_time': round(sum(times) / len(times), 3),
                'min_time': round(min(times), 3),
                'max_time': round(max(times), 3)
            }

    return model_stats

def get_endpoint_stats(self) -> Dict:
    """
    Get endpoint usage statistics.

    Returns:
        dict: Endpoint metrics
    """
    endpoint_stats = {}

    for endpoint, count in self.endpoint_counts.items():
        errors = self.endpoint_errors.get(endpoint, 0)
        error_rate = (errors / count * 100) if count > 0 else 0.0

        endpoint_stats[endpoint] = {
            'requests': count,
            'errors': errors,
            'error_rate': round(error_rate, 2)
        }

    return endpoint_stats

def get_layer_stats(self) -> Dict:
    """
```

Get layer-specific statistics.

Returns:

```
    dict: Layer metrics
"""
stats = {
    'layer1': {
        'anonymity_violations': self.layer_metrics['layer1_anonymity_violations']
    },
    'layer2': {
        'deepfakes_detected': self.layer_metrics['layer2_deepfakes_detected'],
        'avg_credibility_score': round(
            sum(self.layer_metrics['layer2_credibility_scores']) /
            len(self.layer_metrics['layer2_credibility_scores'])
            if self.layer_metrics['layer2_credibility_scores'] else 0.0,
            3
        )
    },
    'layer3': {
        'coordination_detected': self.layer_metrics['layer3_coordination_detected']
    },
    'layer4': {
        'avg_iterations': round(
            sum(self.layer_metrics['layer4_consensus_iterations']) /
            len(self.layer_metrics['layer4_consensus_iterations'])
            if self.layer_metrics['layer4_consensus_iterations'] else 0.0,
            2
        ),
        'avg_convergence_time': round(
            sum(self.layer_metrics['layer4_convergence_times']) /
            len(self.layer_metrics['layer4_convergence_times'])
            if self.layer_metrics['layer4_convergence_times'] else 0.0,
            3
        )
    },
    'layer5': {
        'counter_evidence_count': self.layer_metrics['layer5_counter_evidence_count'],
        'avg_impact_percent': round(
            sum(self.layer_metrics['layer5_impacts']) /
            len(self.layer_metrics['layer5_impacts'])
            if self.layer_metrics['layer5_impacts'] else 0.0,
            2
        )
    },
    'layer6': {
```

```

        'reports_generated': self.layer_metrics['layer6_reports_generated']
    }
}

return stats

def get_security_stats(self) -> Dict:
    """Get security-related statistics."""
    return {
        'hash_chain_failures': self.security_metrics['hash_chain_failures'],
        'validation_failures': dict(self.security_metrics['validation_failures']),
        'crypto_failures': dict(self.security_metrics['crypto_failures']),
        'rate_limit_blocks': dict(self.security_metrics['rate_limit_blocks'])
    }

def get_storage_stats(self) -> Dict:
    """Get storage-related statistics."""
    return dict(self.storage_metrics)

def get_queue_stats(self) -> Dict:
    """Get queue-related statistics."""
    return {
        'pending_jobs': self.queue_metrics['pending_jobs'],
        'processing_jobs': self.queue_metrics['processing_jobs'],
        'job_failures': dict(self.queue_metrics['job_failures'])
    }

def get_evaluation_stats(self) -> Dict:
    """Get research evaluation statistics."""
    return dict(self.evaluation_metrics)

def get_system_metrics(self) -> Dict:
    """
    Get system resource metrics.

    Returns:
        dict: System metrics
    """
    try:
        # Memory info
        memory_info = self.process.memory_info()
        memory_percent = self.process.memory_percent()

        # CPU info
        cpu_percent = self.process.cpu_percent(interval=0.1)
    
```

```

# Uptime
uptime_seconds = time.time() - self.start_time

metrics = {
    'memory_rss_mb': round(memory_info.rss / 1024 / 1024, 2),
    'memory_vms_mb': round(memory_info.vms / 1024 / 1024, 2),
    'memory_percent': round(memory_percent, 2),
    'cpu_percent': round(cpu_percent, 2),
    'num_threads': self.process.num_threads(),
    'uptime_seconds': int(uptime_seconds),
    'uptime_hours': round(uptime_seconds / 3600, 2)
}

return metrics

except Exception as e:
    logger.error(f"Failed to get system metrics: {e}")
    return {}

def get_summary(self) -> Dict:
    """
    Get summary of all metrics.

    Returns:
        dict: Summary metrics for logging/monitoring
    """
    request_stats = self.get_request_stats()
    system_metrics = self.get_system_metrics()
    layer_stats = self.get_layer_stats()

    summary = {
        'requests': request_stats.get('total_requests', 0),
        'error_rate': request_stats.get('error_rate', 0),
        'avg_latency': request_stats.get('avg_latency', 0),
        'memory_percent': system_metrics.get('memory_percent', 0),
        'cpu_percent': system_metrics.get('cpu_percent', 0),
        'uptime_hours': system_metrics.get('uptime_hours', 0),
        'submissions': self.submission_stats.get('total', 0),
        'deepfakes_detected': layer_stats['layer2']['deepfakes_detected'],
        'coordination_detected': layer_stats['layer3']['coordination_detected'],
        'reports_generated': layer_stats['layer6']['reports_generated']
    }

    return summary

def get_all_metrics(self) -> Dict:

```

```

"""
Get all metrics combined.

Returns:
    dict: All metrics
"""

metrics = {
    'timestamp': datetime.utcnow().isoformat(),
    'requests': self.get_request_stats(),
    'throughput_1min': self.get_throughput(60),
    'throughput_5min': self.get_throughput(300),
    'models': self.get_model_stats(),
    'endpoints': self.get_endpoint_stats(),
    'submissions': self.submission_stats.copy(),
    'layers': self.get_layer_stats(),
    'security': self.get_security_stats(),
    'storage': self.get_storage_stats(),
    'queue': self.get_queue_stats(),
    'evaluation': self.get_evaluation_stats(),
    'system': self.get_system_metrics(),
    'generic': self.get_generic_metrics(),
    'prometheus_enabled': self.prometheus_enabled
}

return metrics

def _persist_metrics(self) -> None:
    """Persist metrics to file."""
    if not self.metrics_file:
        return

    try:
        metrics = self.get_all_metrics()

        # Ensure parent directory exists
        self.metrics_file.parent.mkdir(parents=True, exist_ok=True)

        # Append to file (JSON Lines format)
        import json
        with open(self.metrics_file, 'a') as f:
            f.write(json.dumps(metrics) + '\n')

        self.last_persist_time = time.time()

        logger.debug(f"Metrics persisted to {self.metrics_file}")
    
```

```
except Exception as e:
    logger.error(f"Failed to persist metrics: {e}")

def reset_metrics(self) -> None:
    """Reset all metrics."""
    self.request_latencies.clear()
    self.request_timestamps.clear()
    self.model_inference_times.clear()
    self.endpoint_counts.clear()
    self.endpoint_errors.clear()
    self.submission_stats = {
        'total': 0,
        'pending': 0,
        'completed': 0,
        'failed': 0
    }
    self.total_requests = 0
    self.total_errors = 0

    # Reset layer metrics
    self.layer_metrics = {
        'layer1_anonymity_violations': 0,
        'layer2_credibility_scores': deque(maxlen=100),
        'layer2_deepfakes_detected': 0,
        'layer3_coordination_detected': 0,
        'layer4_consensus_iterations': deque(maxlen=100),
        'layer4_convergence_times': deque(maxlen=100),
        'layer5_counter_evidence_count': 0,
        'layer5_impacts': deque(maxlen=100),
        'layer6_reports_generated': 0
    }

    # Reset security metrics
    self.security_metrics = {
        'hash_chain_failures': 0,
        'validation_failures': defaultdict(int),
        'crypto_failures': defaultdict(int),
        'rate_limit_blocks': defaultdict(int)
    }

    # Reset storage metrics
    self.storage_metrics = {
        'corruption_detected': 0,
        'index_inconsistencies': 0,
        'backup_failures': 0
    }
```

```

# Reset queue metrics
self.queue_metrics = {
    'pending_jobs': 0,
    'processing_jobs': 0,
    'job_failures': defaultdict(int)
}

logger.info("Metrics reset")

def get_health_status(self) -> Dict:
    """
    Get system health status.

    Returns:
        dict: Health status
    """
    system_metrics = self.get_system_metrics()
    request_stats = self.get_request_stats()

    # Determine health status
    status = "healthy"
    issues = []

    # Check memory usage
    if system_metrics.get('memory_percent', 0) > 90:
        status = "unhealthy"
        issues.append("High memory usage")
    elif system_metrics.get('memory_percent', 0) > 75:
        status = "degraded"
        issues.append("Elevated memory usage")

    # Check error rate
    error_rate = request_stats.get('error_rate', 0)
    if error_rate > 10:
        status = "unhealthy"
        issues.append("High error rate")
    elif error_rate > 5:
        if status == "healthy":
            status = "degraded"
        issues.append("Elevated error rate")

    health = {
        'status': status,
        'timestamp': datetime.utcnow().isoformat(),
        'uptime_hours': system_metrics.get('uptime_hours', 0),
    }

```

```
        'issues': issues,
        'metrics': {
            'memory_percent': system_metrics.get('memory_percent', 0),
            'cpu_percent': system_metrics.get('cpu_percent', 0),
            'error_rate': error_rate,
            'avg_latency': request_stats.get('avg_latency', 0)
        }
    }

    return health
```

```
def record_gauge(self, metric_name: str, value: float) -> None:
```

```
    """
```

Record a generic gauge metric.

Args:

```
    metric_name: Name of the metric
    value: Metric value
```

```
    """
```

try:

```
    # Store in a generic metrics dict (add to __init__ if needed)
    if not hasattr(self, 'generic_gauges'):
        self.generic_gauges = {}
```

```
    self.generic_gauges[metric_name] = value
```

```
    logger.debug(f"Gauge metric recorded: {metric_name} = {value}")
```

except Exception as e:

```
    logger.error(f"Failed to record gauge metric {metric_name}: {e}")
```

```
def record(self, metric_name: str, value: float) -> None:
```

```
    """
```

Record a generic metric (alias for record_gauge).

Args:

```
    metric_name: Name of the metric
    value: Metric value
```

```
    """
```

```
    self.record_gauge(metric_name, value)
```

```
def get_generic_metrics(self) -> Dict[str, float]:
```

```
    """
```

Get all generic metrics.

Returns:

```
    dict: Generic metrics
    """
    return getattr(self, 'generic_gauges', {})

"""

```

Queue Service - Background job processing

Provides:

- In-memory job queue (no Redis/Celery)
- Async task processing
- Priority queue support
- Job status tracking
- Retry logic with exponential backoff
- Worker pool management

```
"""


```

```
import asyncio
import logging
import time
import uuid
from collections import defaultdict
from dataclasses import dataclass, field
from datetime import datetime
from enum import Enum
from typing import Any, Callable, Dict, List, Optional
```

```
# Initialize logger
logger = logging.getLogger(__name__)
```

```
class JobStatus(str, Enum):
    """Job status enumeration."""
    PENDING = "pending"
    RUNNING = "running"
    COMPLETED = "completed"
    FAILED = "failed"
    RETRYING = "retrying"
    CANCELLED = "cancelled"
```

```
class JobPriority(int, Enum):
    """Job priority levels."""
    LOW = 0
    NORMAL = 1
    HIGH = 2
    CRITICAL = 3
```

```
@dataclass
class Job:
    """
    Job data structure.
    """
    job_id: str
    task_name: str
    func: Callable
    args: tuple = field(default_factory=tuple)
    kwargs: dict = field(default_factory=dict)
    priority: JobPriority = JobPriority.NORMAL
    status: JobStatus = JobStatus.PENDING
    created_at: float = field(default_factory=time.time)
    started_at: Optional[float] = None
    completed_at: Optional[float] = None
    result: Any = None
    error: Optional[str] = None
    retry_count: int = 0
    max_retries: int = 3
    timeout: Optional[int] = None # seconds

    def __lt__(self, other):
        """Compare jobs by priority (for priority queue)."""
        return self.priority.value > other.priority.value # Higher priority first
```

```
class QueueService:
    """
    Queue Service - Background job processing.
    """

    Features:
```

- In-memory priority queue
- Async job processing
- Retry logic with exponential backoff
- Job status tracking
- Worker pool management
- Job history

```
def __init__(
    self,
    max_workers: int = 4,
    max_queue_size: int = 1000,
    job_timeout: int = 300
```

```
):
    """
    Initialize queue service.

    Args:
        max_workers: Maximum concurrent workers
        max_queue_size: Maximum queue size
        job_timeout: Default job timeout in seconds
    """
    self.max_workers = max_workers
    self.max_queue_size = max_queue_size
    self.default_timeout = job_timeout

    # Job storage
    self.jobs: Dict[str, Job] = {}
    self.pending_queue: asyncio.PriorityQueue =
        asyncio.PriorityQueue(maxsize=max_queue_size)

    # Worker management
    self.workers: List[asyncio.Task] = []
    self.running = False
    self.queues = defaultdict(list)

    # Statistics
    self.stats = {
        'jobs_submitted': 0,
        'jobs_completed': 0,
        'jobs_failed': 0,
        'jobs_retried': 0,
        'total_processing_time': 0.0
    }

    logger.info(
        f"QueueService initialized "
        f"(workers={max_workers}, queue_size={max_queue_size})"
    )

async def start(self) -> None:
    """
    Start worker pool.
    """
    if self.running:
        logger.warning("Queue service already running")
        return

    self.running = True

    # Start worker tasks
```

```
for i in range(self.max_workers):
    worker = asyncio.create_task(self._worker(worker_id=i))
    self.workers.append(worker)

logger.info(f"Started {self.max_workers} workers")

async def stop(self) -> None:
    """Stop worker pool gracefully."""
    if not self.running:
        return

    self.running = False

    # Wait for workers to finish
    if self.workers:
        await asyncio.gather(*self.workers, return_exceptions=True)

    self.workers.clear()
    logger.info("Queue service stopped")

async def submit_job(
    self,
    task_name: str,
    func: Callable,
    *args,
    priority: JobPriority = JobPriority.NORMAL,
    max_retries: int = 3,
    timeout: Optional[int] = None,
    **kwargs
) -> str:
    """
    Submit job to queue.

    Args:
        task_name: Human-readable task name
        func: Function to execute
        *args: Positional arguments
        priority: Job priority
        max_retries: Maximum retry attempts
        timeout: Job timeout in seconds
        **kwargs: Keyword arguments
    """

    Submit job to queue.
```

Args:

- task_name: Human-readable task name
- func: Function to execute
- *args: Positional arguments
- priority: Job priority
- max_retries: Maximum retry attempts
- timeout: Job timeout in seconds
- **kwargs: Keyword arguments

Returns:

- str: Job ID

Raises:

```

    ValueError: If queue is full
    """
try:
    # Check queue size
    if self.pending_queue.qsize() >= self.max_queue_size:
        raise ValueError("Queue is full")

    # Create job
    job_id = str(uuid.uuid4())

    job = Job(
        job_id=job_id,
        task_name=task_name,
        func=func,
        args=args,
        kwargs=kwargs,
        priority=priority,
        max_retries=max_retries,
        timeout=timeout or self.default_timeout
    )

    # Store job
    self.jobs[job_id] = job

    # Add to queue
    await self.pending_queue.put(job)

    self.stats['jobs_submitted'] += 1

    logger.info(
        f"Job submitted: {job_id} ({task_name}) "
        f"[priority={priority.name}]"
    )

    return job_id

except Exception as e:
    logger.error(f"Failed to submit job: {e}")
    raise

def enqueue(self, queue_name: str, item: str) -> None:
    """Add item to queue."""
    self.queues[queue_name].append(item)

def dequeue(self, queue_name: str) -> Optional[str]:
    """Remove and return item from queue."""

```

```

if self.queues[queue_name]:
    return self.queues[queue_name].pop(0)
return None

def peek(self, queue_name: str) -> Optional[str]:
    """View next item without removing."""
    if self.queues[queue_name]:
        return self.queues[queue_name][0]
    return None

def size(self, queue_name: str) -> int:
    """Get queue size."""
    return len(self.queues[queue_name])

async def _worker(self, worker_id: int) -> None:
    """
    Worker coroutine that processes jobs from queue.

    Args:
        worker_id: Worker identifier
    """
    logger.info(f"Worker {worker_id} started")

    while self.running:
        try:
            # Get job from queue (with timeout)
            try:
                job = await asyncio.wait_for(
                    self.pending_queue.get(),
                    timeout=1.0
                )
            except asyncio.TimeoutError:
                continue

            # Process job
            await self._process_job(job, worker_id)

        except Exception as e:
            logger.error(f"Worker {worker_id} error: {e}")

    logger.info(f"Worker {worker_id} stopped")

async def _process_job(self, job: Job, worker_id: int) -> None:
    """
    Process a single job.
    """

```

```

Args:
    job: Job to process
    worker_id: Worker identifier
    """
try:
    # Update job status
    job.status = JobStatus.RUNNING
    job.started_at = time.time()

    logger.info(
        f"Worker {worker_id} processing job {job.job_id} "
        f"({job.task_name})"
    )
    # Execute job with timeout
    try:
        if asyncio.iscoroutinefunction(job.func):
            # Async function
            result = await asyncio.wait_for(
                job.func(*job.args, **job.kwargs),
                timeout=job.timeout
            )
        else:
            # Sync function - run in executor
            loop = asyncio.get_event_loop()
            result = await asyncio.wait_for(
                loop.run_in_executor(
                    None,
                    lambda: job.func(*job.args, **job.kwargs)
                ),
                timeout=job.timeout
            )
    except asyncio.TimeoutError:
        logger.error(f"Job {job.job_id} timed out after {job.timeout} seconds")
        job.status = JobStatus.TIMED_OUT
        job.result = None
        job.completed_at = time.time()

    # Job completed successfully
    job.status = JobStatus.COMPLETED
    job.result = result
    job.completed_at = time.time()

    # Update stats
    processing_time = job.completed_at - job.started_at
    self.stats['jobs_completed'] += 1
    self.stats['total_processing_time'] += processing_time

    logger.info(
        f"Job completed: {job.job_id} "
        f"(time={processing_time:.2f}s)"
    )

```

```
        )

    except asyncio.TimeoutError:
        raise TimeoutError(f"Job exceeded timeout of {job.timeout}s")

    except Exception as e:
        # Job failed
        error_msg = str(e)
        logger.error(f"Job failed: {job.job_id} - {error_msg}")

        job.error = error_msg

        # Retry logic
        if job.retry_count < job.max_retries:
            job.retry_count += 1
            job.status = JobStatus.RETRYING

            # Calculate backoff delay (exponential)
            backoff_delay = min(2 ** job.retry_count, 60) # Max 60s

            logger.info(
                f"Retrying job {job.job_id} "
                f"(attempt {job.retry_count}/{job.max_retries}) "
                f"in {backoff_delay}s"
            )

        # Schedule retry
        await asyncio.sleep(backoff_delay)
        await self.pending_queue.put(job)

        self.stats['jobs_retried'] += 1
    else:
        # Max retries exceeded
        job.status = JobStatus.FAILED
        job.completed_at = time.time()

        self.stats['jobs_failed'] += 1

        logger.error(
            f"Job failed permanently: {job.job_id} "
            f"(retries={job.retry_count})"
        )

def get_job_status(self, job_id: str) -> Optional[Dict]:
    """
    Get job status.
    """
```

```

Args:
    job_id: Job identifier

Returns:
    dict: Job status information or None
"""
job = self.jobs.get(job_id)

if not job:
    return None

status_info = {
    'job_id': job.job_id,
    'task_name': job.task_name,
    'status': job.status.value,
    'priority': job.priority.name,
    'created_at': datetime.fromtimestamp(job.created_at).isoformat(),
    'retry_count': job.retry_count,
    'max_retries': job.max_retries
}

if job.started_at:
    status_info['started_at'] = datetime.fromtimestamp(job.started_at).isoformat()

if job.completed_at:
    status_info['completed_at'] =
    datetime.fromtimestamp(job.completed_at).isoformat()
    status_info['duration'] = job.completed_at - job.started_at

if job.status == JobStatus.COMPLETED:
    status_info['result'] = job.result
elif job.status == JobStatus.FAILED:
    status_info['error'] = job.error

return status_info

def get_job_result(self, job_id: str) -> Optional[Any]:
"""
Get job result (if completed).

Args:
    job_id: Job identifier

Returns:
    Job result or None

```

```
"""
job = self.jobs.get(job_id)

if job and job.status == JobStatus.COMPLETED:
    return job.result

return None

async def cancel_job(self, job_id: str) -> bool:
    """
    Cancel pending job.

    Args:
        job_id: Job identifier

    Returns:
        bool: True if cancelled, False if not found or already running
    """
    job = self.jobs.get(job_id)

    if not job:
        return False

    if job.status == JobStatus.PENDING:
        job.status = JobStatus.CANCELLED
        logger.info(f"Job cancelled: {job_id}")
        return True

    return False

def get_queue_stats(self) -> Dict:
    """
    Get queue statistics.

    Returns:
        dict: Queue statistics
    """

    # Count jobs by status
    status_counts = defaultdict(int)
    for job in self.jobs.values():
        status_counts[job.status.value] += 1

    # Calculate average processing time
    avg_processing_time = 0.0
    if self.stats['jobs_completed'] > 0:
        avg_processing_time = (
```

```

        self.stats['total_processing_time'] /
        self.stats['jobs_completed']
    )

stats = {
    'running': self.running,
    'workers': len(self.workers),
    'queue_size': self.pending_queue.qsize(),
    'max_queue_size': self.max_queue_size,
    'total_jobs': len(self.jobs),
    'jobs_by_status': dict(status_counts),
    'jobs_submitted': self.stats['jobs_submitted'],
    'jobs_completed': self.stats['jobs_completed'],
    'jobs_failed': self.stats['jobs_failed'],
    'jobs_retried': self.stats['jobs_retried'],
    'avg_processing_time': round(avg_processing_time, 2),
    'timestamp': datetime.utcnow().isoformat()
}

```

return stats

def cleanup_old_jobs(self, max_age_seconds: int = 86400) -> int:

"""

Remove old completed/failed jobs.

Args:

max_age_seconds: Maximum age in seconds (default: 24 hours)

Returns:

int: Number of jobs removed

"""

current_time = time.time()

cutoff_time = current_time - max_age_seconds

jobs_to_remove = []

for job_id, job in self.jobs.items():

Remove if completed/failed and old

if job.status in (JobStatus.COMPLETED, JobStatus.FAILED,

JobStatus.CANCELLED):

if job.completed_at and job.completed_at < cutoff_time:

jobs_to_remove.append(job_id)

for job_id in jobs_to_remove:

del self.jobs[job_id]

```
    if jobs_to_remove:
        logger.info(f"Cleaned up {len(jobs_to_remove)} old jobs")

    return len(jobs_to_remove)

"""

```

Rate Limiter Service - IP-based rate limiting

Provides:

- IP-based request rate limiting
- Sliding window algorithm
- Configurable limits per endpoint
- Temporary IP blocking
- Rate limit status tracking

```
"""

```

```
import logging
import time
from collections import defaultdict
from datetime import datetime, timedelta
from typing import Dict, Optional, Tuple

# Initialize logger
logger = logging.getLogger(__name__)


```

class RateLimiter:

```
"""

```

Rate Limiter - IP-based request rate limiting.

Features:

- Sliding window rate limiting
- Per-IP tracking
- Configurable limits (requests per time window)
- Automatic cleanup of old records
- Temporary IP blocking for abuse
- Different limits for different endpoint types

```
"""

```

```
def __init__(
    self,
    default_limit: int = 10,
    window_seconds: int = 3600,
    cleanup_interval: int = 300
):
    """

```

Initialize rate limiter.

Args:

```
    default_limit: Default number of requests per window
    window_seconds: Time window in seconds (default: 1 hour)
    cleanup_interval: Cleanup old records every N seconds
```

"""

```
self.default_limit = default_limit
self.window_seconds = window_seconds
self.cleanup_interval = cleanup_interval
```

```
# IP tracking: {ip_address: [timestamps]}
self.requests: Dict[str, list] = defaultdict(list)
```

```
# Blocked IPs: {ip_address: unblock_time}
self.blocked_ips: Dict[str, float] = {}
```

```
# Last cleanup time
self.last_cleanup = time.time()
```

```
# Custom limits for specific endpoint types
```

```
self.endpoint_limits = {
    'submission': 1000, # Increased for evaluation
    'status': 3600, # 1 per second
    'report': 100, # Increased
    'health': 3600, # 1 per second
    'counter_evidence': 100 # Increased
}
```

```
logger.info(
    f"RateLimiter initialized "
    f"(default={default_limit} req/{window_seconds}s)"
)
```

```
def check_rate_limit(
    self,
    ip_address: str,
    endpoint_type: str = 'default'
) -> Tuple[bool, Dict]:
```

"""

Check if request is within rate limit.

Args:

```
    ip_address: Client IP address
    endpoint_type: Type of endpoint (submission/status/report/etc)
```

Returns:
tuple: (is_allowed, rate_limit_info)

"""

```

try:
    current_time = time.time()

    # Cleanup old records periodically
    if current_time - self.last_cleanup > self.cleanup_interval:
        self._cleanup_old_records()

    # Check if IP is blocked
    if ip_address in self.blocked_ips:
        unblock_time = self.blocked_ips[ip_address]

        if current_time < unblock_time:
            remaining = int(unblock_time - current_time)
            logger.warning(
                f"Blocked IP attempted request: {ip_address} "
                f"(blocked for {remaining}s)"
            )

    return False, {
        'allowed': False,
        'reason': 'IP temporarily blocked',
        'retry_after': remaining,
        'blocked_until': datetime.fromtimestamp(unblock_time).isoformat()
    }
else:
    # Unblock IP
    del self.blocked_ips[ip_address]
    logger.info(f"IP unblocked: {ip_address}")

    # Get limit for endpoint type
    limit = self.endpoint_limits.get(endpoint_type, self.default_limit)

    # Get request history for IP
    request_times = self.requests[ip_address]

    # Remove requests outside time window
    cutoff_time = current_time - self.window_seconds
    request_times = [t for t in request_times if t > cutoff_time]
    self.requests[ip_address] = request_times

    # Check if under limit
    current_count = len(request_times)

```

```

if current_count >= limit:
    # Rate limit exceeded
    logger.warning(
        f"Rate limit exceeded: {ip_address} "
        f"({current_count}/{limit} for {endpoint_type})"
    )

    # Check if should block IP (excessive violations)
    if current_count >= limit * 2:
        self._block_ip(ip_address, duration=3600) # Block for 1 hour

    # Calculate when next request is allowed
    oldest_request = min(request_times) if request_times else current_time
    retry_after = int(self.window_seconds - (current_time - oldest_request))

    return False, {
        'allowed': False,
        'reason': 'Rate limit exceeded',
        'limit': limit,
        'current': current_count,
        'window_seconds': self.window_seconds,
        'retry_after': retry_after,
        'endpoint_type': endpoint_type
    }

# Add current request
request_times.append(current_time)

# Calculate remaining requests
remaining = limit - (current_count + 1)
reset_time = current_time + self.window_seconds

logger.debug(
    f"Rate limit check passed: {ip_address} "
    f"({current_count + 1}/{limit} for {endpoint_type})"
)

return True, {
    'allowed': True,
    'limit': limit,
    'remaining': remaining,
    'reset': int(reset_time),
    'reset_datetime': datetime.fromtimestamp(reset_time).isoformat(),
    'window_seconds': self.window_seconds,
    'endpoint_type': endpoint_type
}

```

```
except Exception as e:
    logger.error(f"Rate limit check failed: {e}")
    # Fail open (allow request on error)
    return True, {
        'allowed': True,
        'error': str(e)
    }

def _block_ip(self, ip_address: str, duration: int = 3600) -> None:
    """
    Temporarily block IP address.

    Args:
        ip_address: IP to block
        duration: Block duration in seconds
    """
    unblock_time = time.time() + duration
    self.blocked_ips[ip_address] = unblock_time

    logger.warning(
        f"IP blocked: {ip_address} "
        f"(duration={duration}s,"
        until={datetime.fromtimestamp(unblock_time).isoformat()}"
    )

def unblock_ip(self, ip_address: str) -> bool:
    """
    Manually unblock IP address.

    Args:
        ip_address: IP to unblock

    Returns:
        bool: True if IP was blocked, False if not
    """
    if ip_address in self.blocked_ips:
        del self.blocked_ips[ip_address]
        logger.info(f"IP manually unblocked: {ip_address}")
        return True
    return False

def _cleanup_old_records(self) -> None:
    """
    Remove old request records to free memory.
    """
    try:
        current_time = time.time()
```

```

cutoff_time = current_time - self.window_seconds

# Cleanup request history
cleaned_ips = []
for ip, timestamps in list(self.requests.items()):
    # Remove old timestamps
    timestamps[:] = [t for t in timestamps if t > cutoff_time]

    # Remove IP if no recent requests
    if not timestamps:
        cleaned_ips.append(ip)

for ip in cleaned_ips:
    del self.requests[ip]

# Cleanup expired blocks
expired_blocks = [
    ip for ip, unblock_time in self.blocked_ips.items()
    if current_time >= unblock_time
]

for ip in expired_blocks:
    del self.blocked_ips[ip]

self.last_cleanup = current_time

if cleaned_ips or expired_blocks:
    logger.debug(
        f"Cleanup: removed {len(cleaned_ips)} IPs, "
        f"unblocked {len(expired_blocks)} IPs"
    )

```

except Exception as e:
 logger.error(f"Cleanup failed: {e}")

```

def get_rate_limit_status(self, ip_address: str) -> Dict:
    """

```

Get current rate limit status for IP.

Args:

ip_address: IP address

Returns:

dict: Rate limit status

"""

current_time = time.time()

```

cutoff_time = current_time - self.window_seconds

# Get recent requests
request_times = [t for t in self.requests.get(ip_address, []) if t > cutoff_time]

# Check if blocked
is_blocked = ip_address in self.blocked_ips
unblock_time = self.blocked_ips.get(ip_address, 0)

status = {
    'ip_address': ip_address,
    'is_blocked': is_blocked,
    'recent_requests': len(request_times),
    'limits': self.endpoint_limits,
    'window_seconds': self.window_seconds
}

if is_blocked:
    remaining = int(unblock_time - current_time)
    status['blocked_remaining'] = remaining
    status['unblock_time'] = datetime.fromtimestamp(unblock_time).isoformat()

return status

def reset_ip(self, ip_address: str) -> None:
    """
    Reset rate limit counters for IP.

    Args:
        ip_address: IP address
    """
    if ip_address in self.requests:
        del self.requests[ip_address]

    if ip_address in self.blocked_ips:
        del self.blocked_ips[ip_address]

    logger.info(f"Rate limit reset for IP: {ip_address}")

def get_statistics(self) -> Dict:
    """
    Get rate limiter statistics.

    Returns:
        dict: Statistics
    """

```

```

current_time = time.time()
cutoff_time = current_time - self.window_seconds

# Count active IPs (with recent requests)
active_ips = sum(
    1 for timestamps in self.requests.values()
    if any(t > cutoff_time for t in timestamps)
)

# Count total recent requests
total_requests = sum(
    len([t for t in timestamps if t > cutoff_time])
    for timestamps in self.requests.values()
)

stats = {
    'active_ips': active_ips,
    'blocked_ips': len(self.blocked_ips),
    'total_tracked_ips': len(self.requests),
    'recent_requests': total_requests,
    'window_seconds': self.window_seconds,
    'default_limit': self.default_limit,
    'endpoint_limits': self.endpoint_limits,
    'timestamp': datetime.utcnow().isoformat()
}

```

```
return stats
```

```
def update_limit(self, endpoint_type: str, new_limit: int) -> None:
```

```
"""
```

```
    Update rate limit for endpoint type.
```

```
Args:
```

```
    endpoint_type: Endpoint type
    new_limit: New limit value
"""
```

```
old_limit = self.endpoint_limits.get(endpoint_type, self.default_limit)
self.endpoint_limits[endpoint_type] = new_limit
```

```
logger.info(
    f"Rate limit updated for {endpoint_type}: "
    f"{old_limit} ! {new_limit}"
)
```

```
"""
```

```
Storage Service - JSON-based submission storage
```

Provides:

- Atomic file operations with locking
- Submission CRUD operations
- Validator pool management
- Index management for fast lookup
- Directory sharding for scalability
- Archive functionality

"""

```
import json
import logging
import threading
import os
import time

from datetime import datetime
from pathlib import Path
from typing import Dict, List, Optional
from filelock import FileLock

# Initialize logger
logger = logging.getLogger(__name__)

class StorageService:
    """
    Storage Service for JSON-based data persistence.

    Features:
    - Thread-safe atomic writes with file locking
    - Submission storage with automatic indexing
    - Validator pool persistence
    - Archive functionality
    - Fast lookup via index file
    - Directory sharding by year/month
    """

    def __init__(self, data_dir: Optional[Path] = None):
        """
        Initialize storage service.

        Args:
            data_dir: Base data directory (defaults to backend/data)
        """

        if data_dir:
```

```

        self.data_dir = Path(data_dir)
    else:
        self.data_dir = Path(__file__).parent.parent / "data"

    # Create directory structure
    self.submissions_dir = self.data_dir / "submissions"
    self.evidence_dir = self.data_dir / "evidence"
    self.reports_dir = self.data_dir / "reports"
    self.cache_dir = self.data_dir / "cache"
    self.archive_dir = self.data_dir / "archive"

    # Initialize directories
    self._initialize_directories()

    # Index file
    self.index_file = self.data_dir / "index.json"
    self.index_lock = threading.RLock()

    # Validators file
    self.validators_file = self.data_dir / "validators.json"

    logger.info(f"StorageService initialized (data_dir={self.data_dir})")

def _initialize_directories(self) -> None:
    """Create all required directories."""
    directories = [
        self.data_dir,
        self.submissions_dir,
        self.evidence_dir,
        self.reports_dir,
        self.cache_dir,
        self.archive_dir
    ]

    for directory in directories:
        directory.mkdir(parents=True, exist_ok=True)

    logger.debug(f"Storage directories initialized")

def save_submission(self, submission_id: str, data: Dict) -> None:
    """
    Save submission data to JSON file with atomic write.

    Args:
        submission_id: Unique submission identifier
        data: Submission data dictionary
    """

```

```

Raises:
    IOError: If save fails
"""

try:
    # Ensure submission_id is in data
    data['id'] = submission_id

    # Add/update timestamp
    if 'timestamp_updated' not in data:
        data['timestamp_updated'] = datetime.utcnow().isoformat()

    # File path
    file_path = self.submissions_dir / f"{submission_id}.json"
    lock_path = self.submissions_dir / f"{submission_id}.json.lock"

    # Atomic write with file lock
    lock = FileLock(str(lock_path), timeout=10)

    with lock:
        # Write to temporary file first
        temp_path = file_path.with_suffix('.tmp')

        with open(temp_path, 'w', encoding='utf-8') as f:
            json.dump(data, f, indent=2, ensure_ascii=False)

        # Atomic rename with Windows support
        self._safe_replace(temp_path, file_path)

    # Update index
    self._update_index(submission_id, data)

    logger.debug(f"Saved submission {submission_id}")

except Exception as e:
    logger.error(f"Failed to save submission {submission_id}: {e}")
    raise IOError(f"Failed to save submission: {str(e)}")

def load_submission(self, submission_id: str) -> Optional[Dict]:
    """
    Load submission data from JSON file.

    Args:
        submission_id: Unique submission identifier

    Returns:

```

```

    dict: Submission data, or None if not found
    """
try:
    file_path = self.submissions_dir / f"{submission_id}.json"

    if not file_path.exists():
        logger.debug(f"Submission {submission_id} not found")
        return None

    lock_path = self.submissions_dir / f"{submission_id}.json.lock"
    lock = FileLock(str(lock_path), timeout=10)

    with lock:
        with open(file_path, 'r', encoding='utf-8') as f:
            data = json.load(f)

    return data

except Exception as e:
    logger.error(f"Failed to load submission {submission_id}: {e}")
    return None

def get_submission(self, submission_id: str) -> Optional[Dict]:
    """Alias for load_submission."""
    return self.load_submission(submission_id)

def update_submission(self, submission_id: str, updates: Dict) -> None:
    """
    Update submission with partial data.

    Args:
        submission_id: Submission identifier
        updates: Dictionary of updates to apply
    """
    try:
        lock_path = self.submissions_dir / f"{submission_id}.json.lock"
        lock = FileLock(str(lock_path), timeout=10)

        with lock:
            submission = self.load_submission(submission_id)
            if submission:
                submission.update(updates)
                self.save_submission(submission_id, submission)

    except Exception as e:
        logger.error(f"Failed to update submission {submission_id}: {e}")

```

```
def delete_submission(self, submission_id: str) -> bool:
    """
    Delete submission from storage.

    Args:
        submission_id: Unique submission identifier

    Returns:
        bool: True if deleted, False otherwise
    """
    try:
        file_path = self.submissions_dir / f"{submission_id}.json"

        if not file_path.exists():
            return False

        lock_path = self.submissions_dir / f"{submission_id}.json.lock"
        lock = FileLock(str(lock_path), timeout=10)

        with lock:
            file_path.unlink()

        # Remove lock file
        if lock_path.exists():
            lock_path.unlink()

        # Remove from index
        self._remove_from_index(submission_id)

        logger.info(f"Deleted submission {submission_id}")
        return True

    except Exception as e:
        logger.error(f"Failed to delete submission {submission_id}: {e}")
        return False

def get_all_submissions(self) -> List[Dict]:
    """
    Get all submissions from storage.

    Returns:
        list: List of all submission dictionaries
    """
    submissions = []

    try:
```

```
# Use index for fast lookup
index = self._load_index()

for submission_id in index.keys():
    submission = self.load_submission(submission_id)
    if submission:
        submissions.append(submission)

logger.debug(f"Loaded {len(submissions)} submissions")

except Exception as e:
    logger.error(f"Failed to get all submissions: {e}")

return submissions

def update_submission_status(self, submission_id: str, status: str) -> None:
    """
    Update submission status.

    Args:
        submission_id: Submission identifier
        status: New status value
    """
    submission = self.load_submission(submission_id)

    if submission:
        submission['status'] = status
        submission['status_updated'] = datetime.utcnow().isoformat()
        self.save_submission(submission_id, submission)

def archive_submission(self, submission_id: str) -> bool:
    """
    Archive old submission by moving to archive directory.

    Args:
        submission_id: Submission identifier

    Returns:
        bool: True if archived, False otherwise
    """
    try:
        source = self.submissions_dir / f"{submission_id}.json"

        if not source.exists():
            return False

    except Exception as e:
        logger.error(f"Failed to archive submission {submission_id}: {e}")

    return True
```

```

# Create archive directory if needed
self.archive_dir.mkdir(parents=True, exist_ok=True)

# Move to archive
dest = self.archive_dir / f"{submission_id}.json"

lock_path = self.submissions_dir / f"{submission_id}.json.lock"
lock = FileLock(str(lock_path), timeout=10)

with lock:
    source.rename(dest)

# Clean up lock file
if lock_path.exists():
    lock_path.unlink()

# Remove from index
self._remove_from_index(submission_id)

logger.info(f"Archived submission {submission_id}")
return True

except Exception as e:
    logger.error(f"Failed to archive submission {submission_id}: {e}")
    return False

def save_validators(self, validators: List[Dict]) -> None:
    """
    Save validator pool to storage.

    Args:
        validators: List of validator dictionaries
    """
    try:
        lock_path = self.validators_file.with_suffix('.lock')
        lock = FileLock(str(lock_path), timeout=10)

        with lock:
            temp_path = self.validators_file.with_suffix('.tmp')

            with open(temp_path, 'w', encoding='utf-8') as f:
                json.dump(validators, f, indent=2, ensure_ascii=False)

            temp_path.replace(self.validators_file)

        logger.debug(f"Saved {len(validators)} validators")
    
```

```
except Exception as e:
    logger.error(f"Failed to save validators: {e}")
    raise IOError(f"Failed to save validators: {str(e)}")

def load_validators(self) -> List[Dict]:
    """
    Load validator pool from storage.

    Returns:
        list: List of validator dictionaries
    """
    try:
        if not self.validators_file.exists():
            logger.debug("Validators file not found, returning empty list")
            return []

        lock_path = self.validators_file.with_suffix('.lock')
        lock = FileLock(str(lock_path), timeout=10)

        with lock:
            with open(self.validators_file, 'r', encoding='utf-8') as f:
                validators = json.load(f)

        logger.debug(f"Loaded {len(validators)} validators")
        return validators

    except Exception as e:
        logger.error(f"Failed to load validators: {e}")
        return []

def health_check(self) -> bool:
    """
    Check if storage is accessible and healthy.

    Returns:
        bool: True if healthy, False otherwise
    """
    try:
        # Check if directories exist
        directories_ok = (
            self.data_dir.exists() and
            self.submissions_dir.exists() and
            self.evidence_dir.exists()
        )
    
```

```

if not directories_ok:
    logger.warning("Storage directories not accessible")
    return False

# Test write/read
test_file = self.data_dir / ".health_check"
test_data = {"timestamp": datetime.utcnow().isoformat()}

with open(test_file, 'w') as f:
    json.dump(test_data, f)

with open(test_file, 'r') as f:
    json.load(f)

test_file.unlink()

logger.debug("Storage health check passed")
return True

except Exception as e:
    logger.error(f"Storage health check failed: {e}")
    return False

def _update_index(self, submission_id: str, data: Dict) -> None:
    """
    Update index file with submission metadata.

    Args:
        submission_id: Submission identifier
        data: Submission data
    """
    try:
        with self.index_lock:
            index = self._load_index()

            # Add/update entry
            index[submission_id] = {
                'pseudonym': data.get('pseudonym'),
                'evidence_type': data.get('evidence_type'),
                'status': data.get('status'),
                'timestamp_submission': data.get('timestamp_submission'),
                'timestamp_updated': data.get('timestamp_updated')
            }

            # Save index
            self._save_index(index)
    
```

```
except Exception as e:
    logger.warning(f"Failed to update index: {e}")

def _remove_from_index(self, submission_id: str) -> None:
    """
    Remove submission from index.

    Args:
        submission_id: Submission identifier
    """
    try:
        with self.index_lock:
            index = self._load_index()

            if submission_id in index:
                del index[submission_id]
                self._save_index(index)

    except Exception as e:
        logger.warning(f"Failed to remove from index: {e}")

def _load_index(self) -> Dict:
    """
    Load index file.

    Returns:
        dict: Index data
    """
    try:
        with self.index_lock:
            if not self.index_file.exists():
                return {}

            with open(self.index_file, 'r', encoding='utf-8') as f:
                return json.load(f)

    except Exception as e:
        logger.warning(f"Failed to load index: {e}")
        return {}

def _save_index(self, index: Dict) -> None:
    """
    Save index file.

    Args:

```

```

    index: Index data
    """
try:
    temp_path = self.index_file.with_suffix('.tmp')

    with open(temp_path, 'w', encoding='utf-8') as f:
        json.dump(index, f, indent=2, ensure_ascii=False)

    self._safe_replace(temp_path, self.index_file)

except Exception as e:
    logger.error(f"Failed to save index: {e}")

def get_storage_statistics(self) -> Dict:
    """
    Get storage statistics.

    Returns:
        dict: Storage statistics
    """
    try:
        index = self._load_index()

        # Count by status
        status_counts = {}
        for entry in index.values():
            status = entry.get('status', 'unknown')
            status_counts[status] = status_counts.get(status, 0) + 1

        # Count by evidence type
        type_counts = {}
        for entry in index.values():
            evidence_type = entry.get('evidence_type', 'unknown')
            type_counts[evidence_type] = type_counts.get(evidence_type, 0) + 1

        # Calculate storage size
        total_size = 0
        for file_path in self.submissions_dir.glob("*.json"):
            total_size += file_path.stat().st_size

        stats = {
            'total_submissions': len(index),
            'status_distribution': status_counts,
            'type_distribution': type_counts,
            'storage_size_bytes': total_size,
            'storage_size_mb': round(total_size / (1024 * 1024), 2),
        }
    
```

```
        'data_directory': str(self.data_dir),
        'timestamp': datetime.utcnow().isoformat()
    }

    return stats
```

```
except Exception as e:
    logger.error(f"Failed to get storage statistics: {e}")
    return {
        'error': str(e),
        'timestamp': datetime.utcnow().isoformat()
    }
```

```
def get_evidence_path(
    self,
    submission_id: str,
    filename: str,
    create_dirs: bool = True
) -> Path:
    """
```

Get path for evidence file with year/month sharding.

Args:

```
    submission_id: Submission identifier
    filename: Evidence filename
    create_dirs: Create directories if they don't exist
```

Returns:

```
    Path: Full path for evidence file
    """
```

```
now = datetime.utcnow()
year = now.strftime('%Y')
month = now.strftime('%m')
```

```
# Sharded directory structure
evidence_path = self.evidence_dir / year / month
```

```
if create_dirs:
    evidence_path.mkdir(parents=True, exist_ok=True)
```

```
return evidence_path / filename
```

```
def cleanup_cache(self, max_age_hours: int = 24) -> int:
    """
```

Cleanup old cache files.

Args:

 max_age_hours: Maximum age of cache files in hours

Returns:

 int: Number of files deleted

"""

try:

 from datetime import timedelta

```
    cutoff_time = datetime.utcnow() - timedelta(hours=max_age_hours)
    deleted = 0
```

```
    for file_path in self.cache_dir.glob("*"):
```

```
        if file_path.is_file():
```

```
            file_time = datetime.fromtimestamp(file_path.stat().st_mtime)
```

```
            if file_time < cutoff_time:
```

```
                file_path.unlink()
```

```
                deleted += 1
```

```
    logger.info(f"Cleaned up {deleted} cache files")
```

```
    return deleted
```

except Exception as e:

```
    logger.error(f"Cache cleanup failed: {e}")
```

```
    return 0
```

def save_evidence_file(

 self,

 submission_id: str,

 filename: str,

 content: bytes,

 evidence_type: str

) -> Path:

"""

Save evidence file content to storage.

Args:

 submission_id: Submission identifier

 filename: Original filename

 content: File content bytes

 evidence_type: Type of evidence

Returns:

 Path: Path to saved file

"""

```

try:
    # Get evidence path with sharding
    safe_filename = self._sanitize_filename(filename)
    file_path = self.get_evidence_path(submission_id, safe_filename)

    # Ensure directory exists
    file_path.parent.mkdir(parents=True, exist_ok=True)

    # Write file content
    with open(file_path, 'wb') as f:
        f.write(content)

    logger.debug(f"Saved evidence file: {file_path}")
    return file_path

except Exception as e:
    logger.error(f"Failed to save evidence file: {e}")
    raise IOError(f"Failed to save evidence file: {str(e)}")

def _sanitize_filename(self, filename: str) -> str:
    """Sanitize filename for safe storage."""
    import re
    # Remove path components
    filename = Path(filename).name
    # Replace unsafe characters
    safe_name = re.sub(r'[^a-zA-Z0-9._-]', '_', filename)
    # Limit length
    if len(safe_name) > 255:
        name, ext = safe_name.rsplit('.', 1) if '.' in safe_name else (safe_name, '')
        safe_name = name[:250] + ('.' + ext if ext else '')
    return safe_name or 'unnamed_file'

def _safe_replace(self, src: Path, dst: Path, retries: int = 3) -> None:
    """
    Safely replace file, handling Windows locking issues.
    """

    Args:
        src: Source path
        dst: Destination path
        retries: Number of retries on PermissionError
    """

    for i in range(retries):
        try:
            src.replace(dst)
            return
        except PermissionError:

```

```
    if i == retries - 1:
        raise

    # On Windows, destination might be open or locked temporarily
    # Try to unlink if it exists
    try:
        if dst.exists():
            dst.unlink()
    except Exception:
        pass # Ignore unlink errors, retry replace

    time.sleep(0.1)
```

```
def save_evidence_file(
```

```
    self,
    submission_id: str,
    filename: str,
    content: bytes,
    evidence_type: str
) -> Path:
```

```
"""
```

```
    Save evidence file content to storage.
```

```
Args:
```

```
    submission_id: Submission identifier
    filename: Original filename
    content: File content bytes
    evidence_type: Type of evidence
```

```
Returns:
```

```
    Path: Path to saved file
```

```
"""
```

```
try:
```

```
    # Get evidence path with sharding
    safe_filename = self._sanitize_filename(filename)
    file_path = self.get_evidence_path(submission_id, safe_filename)
```

```
    # Ensure directory exists
    file_path.parent.mkdir(parents=True, exist_ok=True)
```

```
    # Write file content
```

```
    with open(file_path, 'wb') as f:
        f.write(content)
```

```
    logger.debug(f"Saved evidence file: {file_path}")
    return file_path
```

```

except Exception as e:
    logger.error(f"Failed to save evidence file: {e}")
    raise IOError(f"Failed to save evidence file: {str(e)}")

def _sanitize_filename(self, filename: str) -> str:
    """Sanitize filename for safe storage."""
    import re
    # Remove path components
    filename = Path(filename).name
    # Replace unsafe characters
    safe_name = re.sub(r'[^a-zA-Z0-9._-]', '_', filename)
    # Limit length
    if len(safe_name) > 255:
        name, ext = safe_name.rsplit('.', 1) if '.' in safe_name else (safe_name, '')
        safe_name = name[:250] + ('.' + ext if ext else '')
    return safe_name or 'unnamed_file'

```

"""

Validation Service - Input sanitization and validation

Provides:

- File upload validation (type, size, format)
- Text input sanitization
- Evidence type verification
- Malicious content detection
- Input length limits
- Security checks

"""

```

import logging
import mimetypes
import re
from pathlib import Path
from typing import Dict, List, Optional, Tuple

from PIL import Image

# Initialize logger
logger = logging.getLogger(__name__)

```

class ValidationService:

"""

Validation Service - Input sanitization and security validation.

Features:

- File type and size validation
- Text sanitization (XSS, SQL injection prevention)
- Evidence format verification
- Malicious content detection
- Input length limits
- Filename sanitization

"""

File size limits (bytes)

```
MAX_IMAGE_SIZE = 50 * 1024 * 1024 # 50 MB
MAX_VIDEO_SIZE = 500 * 1024 * 1024 # 500 MB
MAX_AUDIO_SIZE = 100 * 1024 * 1024 # 100 MB
MAX_DOCUMENT_SIZE = 20 * 1024 * 1024 # 20 MB
```

Text limits

```
MAX_TEXT_LENGTH = 10000 # 10k characters
MAX_NARRATIVE_LENGTH = 5000 # 5k characters
```

Allowed file extensions

```
ALLOWED_IMAGE_EXTENSIONS = {'jpg', 'jpeg', 'png', 'gif', 'bmp', 'tiff', 'webp'}
ALLOWED_VIDEO_EXTENSIONS = {'mp4', 'avi', 'mov', 'mkv', 'webm'}
ALLOWED_AUDIO_EXTENSIONS = {'mp3', 'wav', 'm4a', 'ogg', 'flac'}
ALLOWED_DOCUMENT_EXTENSIONS = {'pdf', 'doc', 'docx', 'txt'}
```

MIME type mappings

```
ALLOWED_MIME_TYPES = {
    'image': {'image/jpeg', 'image/png', 'image/gif', 'image/bmp', 'image/tiff', 'image/webp'},
    'video': {'video/mp4', 'video/x-msvideo', 'video/quicktime', 'video/x-matroska', 'video/webm'},
    'audio': {'audio/mpeg', 'audio/wav', 'audio/mp4', 'audio/ogg', 'audio/flac'},
    'document': {'application/pdf', 'application/msword', 'text/plain'}
}
```

def __init__(self):

"""Initialize validation service."""

```
logger.info("ValidationService initialized")
```

def validate_file_content(

```
    self,
    filename: str,
    content: bytes,
    evidence_type: str
) -> Tuple[bool, Optional[str]]:
```

"""

Validate file content (in-memory) for security and format compliance.

Args:

filename: Original filename
content: File content bytes
evidence_type: Type of evidence (image/video/audio/document)

Returns:

tuple: (is_valid, error_message)

"""

try:

Check file size

file_size = len(content)

if evidence_type == 'image':

if file_size > self.MAX_IMAGE_SIZE:

return False, f"Image size exceeds limit ({self.MAX_IMAGE_SIZE // (1024*1024)} MB)"

elif evidence_type == 'video':

if file_size > self.MAX_VIDEO_SIZE:

return False, f"Video size exceeds limit ({self.MAX_VIDEO_SIZE // (1024*1024)} MB)"

elif evidence_type == 'audio':

if file_size > self.MAX_AUDIO_SIZE:

return False, f"Audio size exceeds limit ({self.MAX_AUDIO_SIZE // (1024*1024)} MB)"

elif evidence_type == 'document':

if file_size > self.MAX_DOCUMENT_SIZE:

return False, f"Document size exceeds limit ({self.MAX_DOCUMENT_SIZE // (1024*1024)} MB)"

else:

return False, f"Invalid evidence type: {evidence_type}"

Check file extension

extension = Path(filename).suffix.lower() if filename else "

allowed_extensions = self.get_allowed_extensions(evidence_type)

if extension not in allowed_extensions:

return False, f"Invalid {evidence_type} extension: {extension}"

Verify MIME type / Magic Bytes

For in-memory, we can check magic bytes directly

is_mime_valid, mime_error = self._verify_content_type(content, evidence_type)

if not is_mime_valid:

return False, mime_error

```

# Additional checks for images
if evidence_type == 'image':
    is_image_valid, image_error = self._verify_image_content(content)
    if not is_image_valid:
        return False, image_error

    logger.debug(f"File content validation passed: {filename} ({evidence_type})")
    return True, None

except Exception as e:
    logger.error(f"File content validation failed: {e}")
    return False, f"Validation error: {str(e)}"

def validate_file_upload(
    self,
    file_path: Path,
    evidence_type: str
) -> Tuple[bool, Optional[str]]:
    """
    Validate file upload for security and format compliance.

```

Args:

 file_path: Path to uploaded file
 evidence_type: Type of evidence (image/video/audio/document)

Returns:

 tuple: (is_valid, error_message)
 """

try:

 # Check file exists
 if not file_path.exists():
 return False, "File does not exist"

 # Check file size
 file_size = file_path.stat().st_size

 if evidence_type == 'image':
 if file_size > self.MAX_IMAGE_SIZE:

return False, f"Image size exceeds limit ({self.MAX_IMAGE_SIZE // (1024*1024)} MB)"

 elif evidence_type == 'video':
 if file_size > self.MAX_VIDEO_SIZE:

return False, f"Video size exceeds limit ({self.MAX_VIDEO_SIZE // (1024*1024)} MB)"

 elif evidence_type == 'audio':
 if file_size > self.MAX_AUDIO_SIZE:

```

        return False, f"Audio size exceeds limit ({self.MAX_AUDIO_SIZE // (1024*1024)} MB)"
    elif evidence_type == 'document':
        if file_size > self.MAX_DOCUMENT_SIZE:
            return False, f"Document size exceeds limit ({self.MAX_DOCUMENT_SIZE // (1024*1024)} MB)"
    else:
        return False, f"Invalid evidence type: {evidence_type}"

    # Check file extension
    extension = file_path.suffix.lower()

    allowed_extensions = self.get_allowed_extensions(evidence_type)
    if extension not in allowed_extensions:
        return False, f"Invalid {evidence_type} extension: {extension}"

    # Verify MIME type
    is_mime_valid, mime_error = self._verify_mime_type(file_path, evidence_type)
    if not is_mime_valid:
        return False, mime_error

    # Additional checks for images
    if evidence_type == 'image':
        is_image_valid, image_error = self._verify_image_integrity(file_path)
        if not is_image_valid:
            return False, image_error

    logger.debug(f"File validation passed: {file_path.name} ({evidence_type})")
    return True, None

except Exception as e:
    logger.error(f"File validation failed: {e}")
    return False, f"Validation error: {str(e)}"

def _verify_content_type(self, content: bytes, evidence_type: str) -> Tuple[bool, Optional[str]]:
    """Verify content type from magic bytes."""
    try:
        import io
        # Use magic bytes check similar to _detect_mime_from_header but on content
        header = content[:16]
        mime_type = self._detect_mime_from_bytes(header)

        if not mime_type:
            # If magic bytes fail, we might trust the caller or return generic valid
            # For now, let's be strict only if we can identify.

```

```

        # Actually, if we can't identify, return True to avoid blocking valid files we just
        don't know
        return True, None

    allowed_mimes = self.ALLOWED_MIME_TYPES.get(evidence_type, set())
    if mime_type not in allowed_mimes:
        return False, f"File type mismatch (detected {mime_type})"

    return True, None
except Exception as e:
    logger.warning(f"Content type verification failed: {e}")
    return True, None

def _verify_image_content(self, content: bytes) -> Tuple[bool, Optional[str]]:
    """Verify image content integrity."""
    try:
        import io
        with Image.open(io.BytesIO(content)) as img:
            img.verify()
            if img.size[0] > 10000 or img.size[1] > 10000:
                return False, "Image dimensions too large"
            if img.size[0] < 10 or img.size[1] < 10:
                return False, "Image dimensions too small"
        return True, None
    except Exception as e:
        return False, f"Invalid image content: {e}"

def _verify_mime_type(
    self,
    file_path: Path,
    evidence_type: str
) -> Tuple[bool, Optional[str]]:
    """
    Verify file MIME type matches declared type.
    """

```

Args:

file_path: Path to file
evidence_type: Declared evidence type

Returns:

tuple: (is_valid, error_message)

"""

try:

Guess MIME type from file
mime_type, _ = mimetypes.guess_type(str(file_path))

```

if not mime_type:
    # Try reading file header for common types
    mime_type = self._detect_mime_from_header(file_path)

if not mime_type:
    return False, "Could not determine file type"

# Check if MIME type matches evidence type
allowed_mimes = self.ALLOWED_MIME_TYPES.get(evidence_type, set())

if mime_type not in allowed_mimes:
    return False, f"File type mismatch: expected {evidence_type}, got
{mime_type}"

return True, None

except Exception as e:
    logger.warning(f"MIME type verification failed: {e}")
    return True, None # Don't fail validation on MIME check errors

def _detect_mime_from_bytes(self, header: bytes) -> Optional[str]:
    """Detect MIME type from header bytes."""
    try:
        if header.startswith(b'\xff\xd8\xff'):
            return 'image/jpeg'
        elif header.startswith(b'\x89PNG\r\n\x1a\n'):
            return 'image/png'
        elif header.startswith(b'GIF87a') or header.startswith(b'GIF89a'):
            return 'image/gif'
        elif header.startswith(b'RIFF') and b'WEBP' in header:
            return 'image/webp'
        elif header.startswith(b'%PDF'):
            return 'application/pdf'
        elif header.startswith(b'ID3') or header.startswith(b'\xff\xfb'):
            return 'audio/mpeg'
        return None
    except Exception:
        return None

def _detect_mime_from_header(self, file_path: Path) -> Optional[str]:
    """
    Detect MIME type from file header (magic bytes).
    """

```

Args:

file_path: Path to file

```

Returns:
    str: MIME type or None
    """
try:
    with open(file_path, 'rb') as f:
        header = f.read(16)

    return self._detect_mime_from_bytes(header)

except Exception as e:
    logger.warning(f"Magic byte detection failed: {e}")
    return None

def _verify_image_integrity(self, file_path: Path) -> Tuple[bool, Optional[str]]:
    """
    Verify image can be opened and is not corrupted.

Args:
    file_path: Path to image

Returns:
    tuple: (is_valid, error_message)
    """
try:
    with Image.open(file_path) as img:
        # Try to load image
        img.verify()

        # Check dimensions (reasonable limits)
        if img.size[0] > 10000 or img.size[1] > 10000:
            return False, "Image dimensions too large (max 10000x10000)"

        if img.size[0] < 10 or img.size[1] < 10:
            return False, "Image dimensions too small (min 10x10)"

    return True, None

except Exception as e:
    logger.warning(f"Image integrity check failed: {e}")
    return False, f"Invalid or corrupted image: {str(e)}"

def sanitize_text(self, text: str, max_length: Optional[int] = None) -> str:
    """
    Sanitize text input to prevent XSS and injection attacks.

Args:

```

```
text: Input text
max_length: Maximum allowed length

Returns:
str: Sanitized text
"""
if not text:
    return ""

# Convert to string if not already
text = str(text)

# Apply length limit
if max_length:
    text = text[:max_length]
else:
    text = text[:self.MAX_TEXT_LENGTH]

# Remove null bytes
text = text.replace('\x00', "")

# Remove control characters (except newlines and tabs)
text = ''.join(char for char in text if char.isprintable() or char in '\n\r\t')

# Trim whitespace
text = text.strip()

return text

def validate_text_narrative(self, narrative: str) -> Tuple[bool, Optional[str]]:
"""
Validate text narrative for submission.

Args:
narrative: Text narrative

Returns:
tuple: (is_valid, error_message)
"""
if not narrative:
    return False, "Narrative cannot be empty"

# Check length
if len(narrative) < 10:
    return False, "Narrative too short (minimum 10 characters)"
```

```

if len(narrative) > self.MAX_NARRATIVE_LENGTH:
    return False, f"Narrative too long (maximum {self.MAX_NARRATIVE_LENGTH} characters)"

# Check for suspicious patterns
suspicious_patterns = [
    r'<script', # XSS
    r'javascript:', # XSS
    r'on\w+\s*=', # Event handlers
    r'eval\s*\()', # Code execution
    r'exec\s*\()', # Code execution
]
narrative_lower = narrative.lower()
for pattern in suspicious_patterns:
    if re.search(pattern, narrative_lower):
        logger.warning(f"Suspicious pattern detected in narrative: {pattern}")
        return False, "Narrative contains potentially malicious content"

return True, None

```

```
def sanitize_filename(self, filename: str) -> str:
```

"""
 Sanitize filename to prevent path traversal and other attacks.

Args:

filename: Original filename

Returns:

str: Sanitized filename

"""

if not filename:

return "unnamed_file"

Get base name (remove any path components)

filename = Path(filename).name

Remove or replace dangerous characters

Allow: alphanumeric, dash, underscore, dot

safe_chars = re.sub(r'[^a-zA-Z0-9._-]', '_', filename)

Prevent double extensions (.php.jpg)

safe_chars = safe_chars.replace('..', '_')

Limit length

if len(safe_chars) > 255:

```

# Preserve extension
name, ext = safe_chars.rsplit('.', 1) if '.' in safe_chars else (safe_chars, '')
safe_chars = name[:250] + ('.' + ext if ext else '')

# Ensure not empty
if not safe_chars or safe_chars == '.':
    safe_chars = "unnamed_file"

return safe_chars

def validate_submission_id(self, submission_id: str) -> Tuple[bool, Optional[str]]:
    """
    Validate submission ID format.

    Args:
        submission_id: Submission identifier

    Returns:
        tuple: (is_valid, error_message)
    """
    if not submission_id:
        return False, "Submission ID cannot be empty"

    # Check length
    if len(submission_id) < 8 or len(submission_id) > 64:
        return False, "Invalid submission ID length"

    # Check format (alphanumeric, dash, underscore only)
    if not re.match(r'^[a-zA-Z0-9_-]+$', submission_id):
        return False, "Submission ID contains invalid characters"

    return True, None

def validate_pseudonym(self, pseudonym: str) -> Tuple[bool, Optional[str]]:
    """
    Validate pseudonym format.

    Args:
        pseudonym: Pseudonym string

    Returns:
        tuple: (is_valid, error_message)
    """
    if not pseudonym:
        return False, "Pseudonym cannot be empty"

```

```

# Check length
if len(pseudonym) < 5 or len(pseudonym) > 100:
    return False, "Invalid pseudonym length"

# Should start with expected prefix
if not pseudonym.startswith('whistleblower-'):
    return False, "Invalid pseudonym format"

return True, None

def check_sql_injection(self, text: str) -> bool:
    """
    Check text for SQL injection patterns.

    Args:
        text: Text to check

    Returns:
        bool: True if suspicious, False if clean
    """
    if not text:
        return False

    text_lower = text.lower()

    # Common SQL injection patterns
    sql_patterns = [
        r'union\s+select',
        r';\s*drop\s+table',
        r';\s*delete\s+from',
        r';\s*update\s+',
        r';\s*insert\s+into',
        r'--', # SQL comment
        r'/*.*/*', # SQL comment
        r'xp_cmdshell',
        r'exec\s*\(',
    ]

    for pattern in sql_patterns:
        if re.search(pattern, text_lower):
            logger.warning(f"SQL injection pattern detected: {pattern}")
            return True

    return False

def validate_evidence_type(self, evidence_type: str) -> Tuple[bool, Optional[str]]:

```

```
"""
Validate evidence type.

Args:
    evidence_type: Evidence type string

Returns:
    tuple: (is_valid, error_message)
"""
valid_types = {'image', 'video', 'audio', 'document'}

if not evidence_type:
    return False, "Evidence type cannot be empty"

if evidence_type.lower() not in valid_types:
    return False, f"Invalid evidence type. Must be one of: {', '.join(valid_types)}"

return True, None
```

```
def get_file_size_limit(self, evidence_type: str) -> int:
```

```
"""
Get file size limit for evidence type.

Args:
    evidence_type: Evidence type

Returns:
    int: Size limit in bytes
"""
limits = {
```

```
    'image': self.MAX_IMAGE_SIZE,
    'video': self.MAX_VIDEO_SIZE,
    'audio': self.MAX_AUDIO_SIZE,
    'document': self.MAX_DOCUMENT_SIZE
```

```
}
```

```
return limits.get(evidence_type.lower(), self.MAX_DOCUMENT_SIZE)
```

```
def get_allowed_extensions(self, evidence_type: str) -> set:
```

```
"""
Get allowed file extensions for evidence type.

Args:
    evidence_type: Evidence type

Returns:
```

```
    set: Allowed extensions
    """
extensions = {
    'image': self.ALLOWED_IMAGE_EXTENSIONS,
    'video': self.ALLOWED_VIDEO_EXTENSIONS,
    'audio': self.ALLOWED_AUDIO_EXTENSIONS,
    'document': self.ALLOWED_DOCUMENT_EXTENSIONS
}

return extensions.get(evidence_type.lower(), set())
"""

```

Services Package - Centralized service exports

This module provides centralized imports for all service classes, making it easier to import services throughout the application.

Usage:

```
from backend.services import StorageService, CryptoService
"""

```

```
from backend.services.storage_service import StorageService
from backend.services.hash_chain_service import HashChainService
from backend.services.crypto_service import CryptoService
from backend.services.metadata_service import MetadataService
from backend.services.validation_service import ValidationService
from backend.services.rate_limiter import RateLimiter
from backend.services.queue_service import QueueService, Job, JobStatus, JobPriority
from backend.services.metrics_service import MetricsService
```

```
__all__ = [
    # Storage and data management
    'StorageService',
    'HashChainService',

    # Security services
    'CryptoService',
    'MetadataService',
    'ValidationService',
    'RateLimiter',

    # Background processing
    'QueueService',
    'Job',
    'JobStatus',
    'JobPriority',
```

```
# Monitoring
'MetricsService',
]
__version__ = '0.1.0'
```